

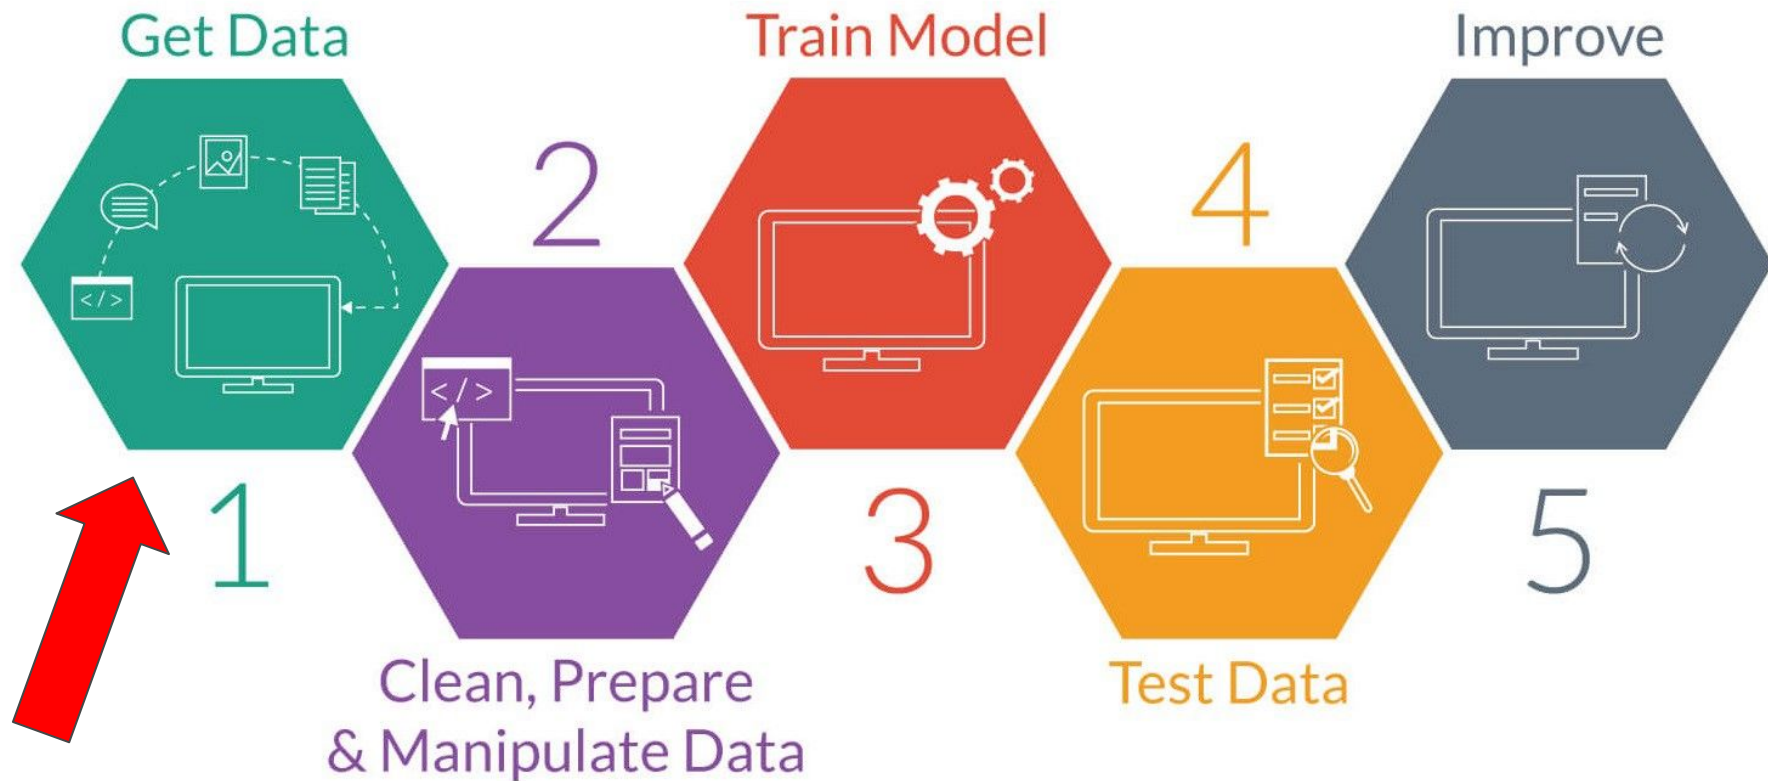


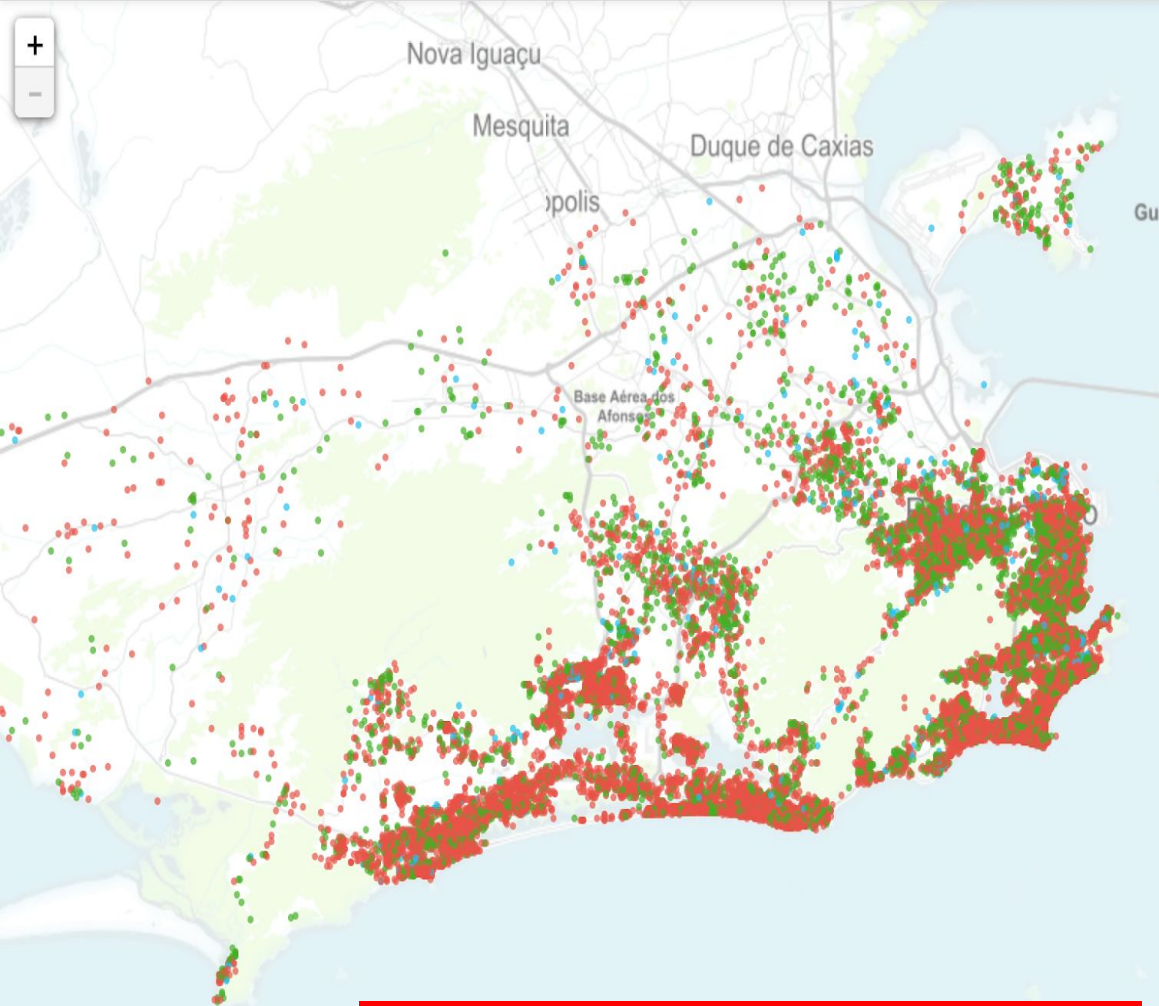
Lesson #03

K-Nearest Neighbors

- Univariate KNN
 - Euclidean distance for univariate
 - Function to make predictions
 - Error metrics
- Multivariate KNN
 - Normalize columns
 - Euclidean distance for multivariate
- Hyperparameter optimization
- Cross-Validation
- Pipeline & Gridsearch

A general ML workflow





Determining the optimal nightly rent price

Rio de Janeiro

Filter by:

Rio de Janeiro

35,887

out of 35,887 listings (100%)

[About Airbnb in Rio de Janeiro](#)


How is Airbnb really being used in and affecting your neighbourhoods?

Room Type

☐ Only entire homes/apartments

Airbnb hosts can list entire homes/apartments, private or shared rooms.

Depending on the room type and [activity](#), an airbnb listing could be more like a hotel, disruptive for neighbours, taking away housing, and [illegal](#).



71.4%
entire homes/apartments

R\$626
price/night

25,629 (71.4%)
entire home/apartments

9,440 (26.3%)
private rooms

818 (2.3%)
shared rooms

- **host_response_rate**: the response rate of the host
- **host_acceptance_rate**: number of requests to the host that convert to rentals
- **host_listings_count**: number of other listings the host has
- **latitude**: latitude dimension of the geographic coordinates
- **longitude**: longitude part of the coordinates
- **city**: the city the living space resides
- **zipcode**: the zip code the living space resides
- **state**: the state the living space resides
- **accommodates**: the number of guests the rental can accommodate
- **room_type**: the type of living space (Private room, Shared room or Entire home/apt)
- **bedrooms**: number of bedrooms included in the rental
- **bathrooms**: number of bathrooms included in the rental
- **beds**: number of beds included in the rental
- **price**: nightly price for the rental
- **cleaning_fee**: additional fee used for cleaning the living space after the guest leaves
- **security_deposit**: refundable security deposit, in case of damages
- **minimum_nights**: minimum number of nights a guest can stay for the rental
- **maximum_nightss**: maximum number of nights a guest can stay for the rental
- **number_of_reviews**: number of reviews that previous guests have left


106
cols

Select the number of similar listings, k , you want to compare with.


$k = 3$



$k = 5$



$k = 7$



For this example, we'll use 3 for our k value.

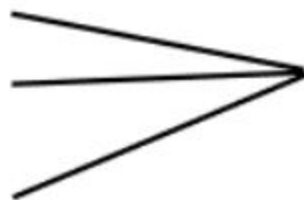
dataset

bedrooms	price
1	160
3	350
1	60
1	95
1	50

Rank each listing by the similarity metric and select the first **k** listings.

dataset (ordered
by similarity)

bedrooms	price	similarity
1	160	0
1	60	0
1	95	0
1	50	0
3	350	2



mean price

105

Euclidean distance - Univariate

Univariate case

$$d = \sqrt{(q_1 - p_1)^2}$$

$$d = |q_1 - p_1|$$

rio_listings

index	accommodates	
	our listing	
		8
index	accommodates	distance
0	4	$(4 - 8)^2$
1	6	$(6 - 8)^2$
2	1	$(1 - 8)^2$
3	2	$(2 - 8)^2$

Euclidean distance (multivariate)

$$d = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \dots + (q_n - p_n)^2}$$

index	host_listings_count	accommodates	bedrooms	bathrooms	beds
0	26	4	1	1	2
1	1	6	3	3	3

$$(q_1 - p_1) + (q_2 - p_2) + \dots + (q_n - p_n) \quad \text{differences} \quad (26 - 1) + (4 - 6) + (1 - 3) + (1 - 3) + (2 - 3)$$

$$(q_1 - p_1)^2 + (q_2 - p_2)^2 + \dots + (q_n - p_n)^2 \quad \text{squared differences} \quad (25)^2 + (-2)^2 + (-2)^2 + (-2)^2 + (-1)^2$$

$$\sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \dots + (q_n - p_n)^2}$$

Euclidean distance	= $\sqrt{625 + 4 + 4 + 4 + 1}$
	= $\sqrt{638}$
	= 25.258661

Error metrics (regression problem)

Mean Absolute Error

$$MAE = \frac{|actual_1 - predicted_1| + |actual_2 - predicted_2| + \dots + |actual_n - predicted_n|}{n}$$

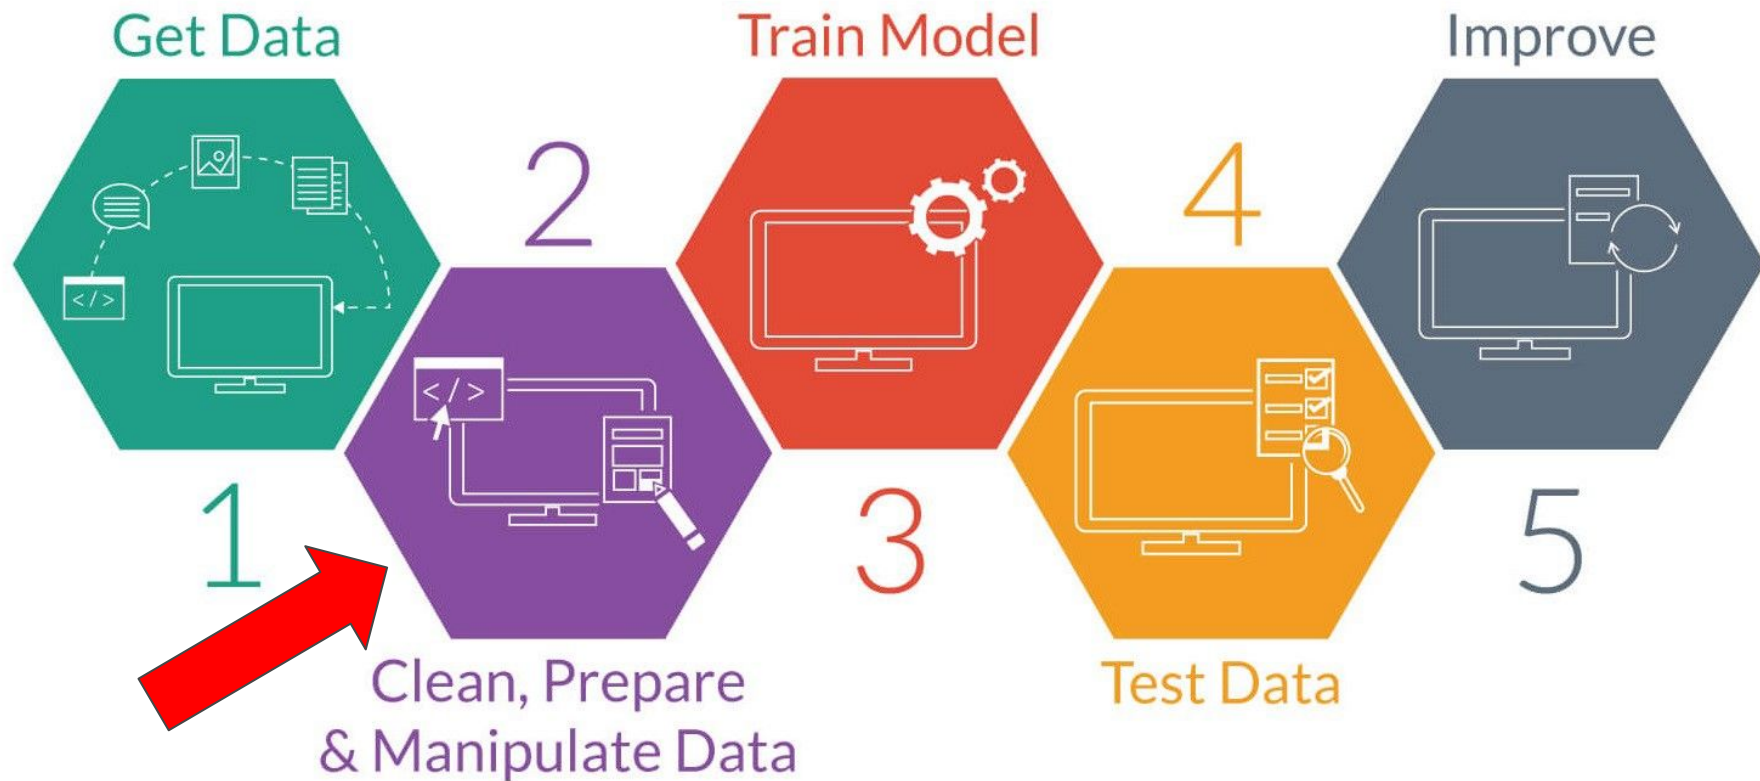
Mean Squared Error

$$MSE = \frac{(actual_1 - predicted_1)^2 + (actual_2 - predicted_2)^2 + \dots + (actual_n - predicted_n)^2}{n}$$

Root Mean Squared Error

$$RMSE = \sqrt{MSE}$$

A general ML workflow



Removing features


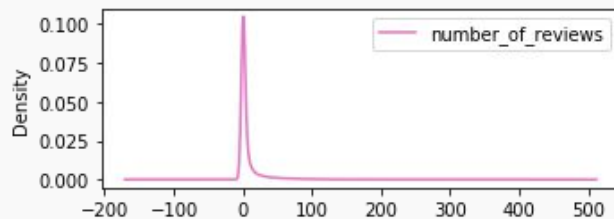

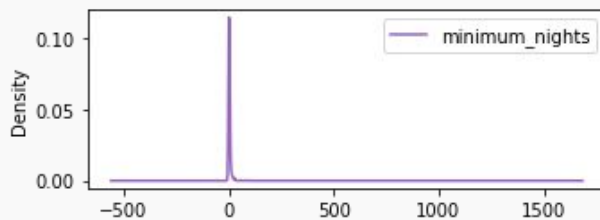



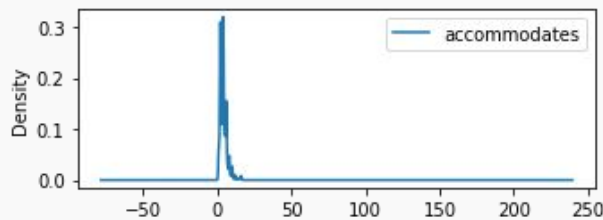
- `room_type`: e.g. **Private room**
- `city`: e.g. **Washington**
- `state`: e.g. **DC**
- `host_response_rate`
- `host_acceptance_rate`
- `host_listings_count`
- `latitude`: e.g. **38.913458**
- `longitude`: e.g. **-77.031**
- `zipcode`: e.g. **20009**

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 35736 entries, 30160 to 33003
Data columns (total 8 columns):
accommodates      35736 non-null int64
bedrooms          35714 non-null float64
bathrooms         35660 non-null float64
beds              35693 non-null float64
price             35736 non-null float64
cleaning_fee      23389 non-null object
minimum_nights    35736 non-null int64
maximum_nights    35736 non-null int64
dtypes: float64(4), int64(3), object(1)
memory usage: 2.5+ MB
```

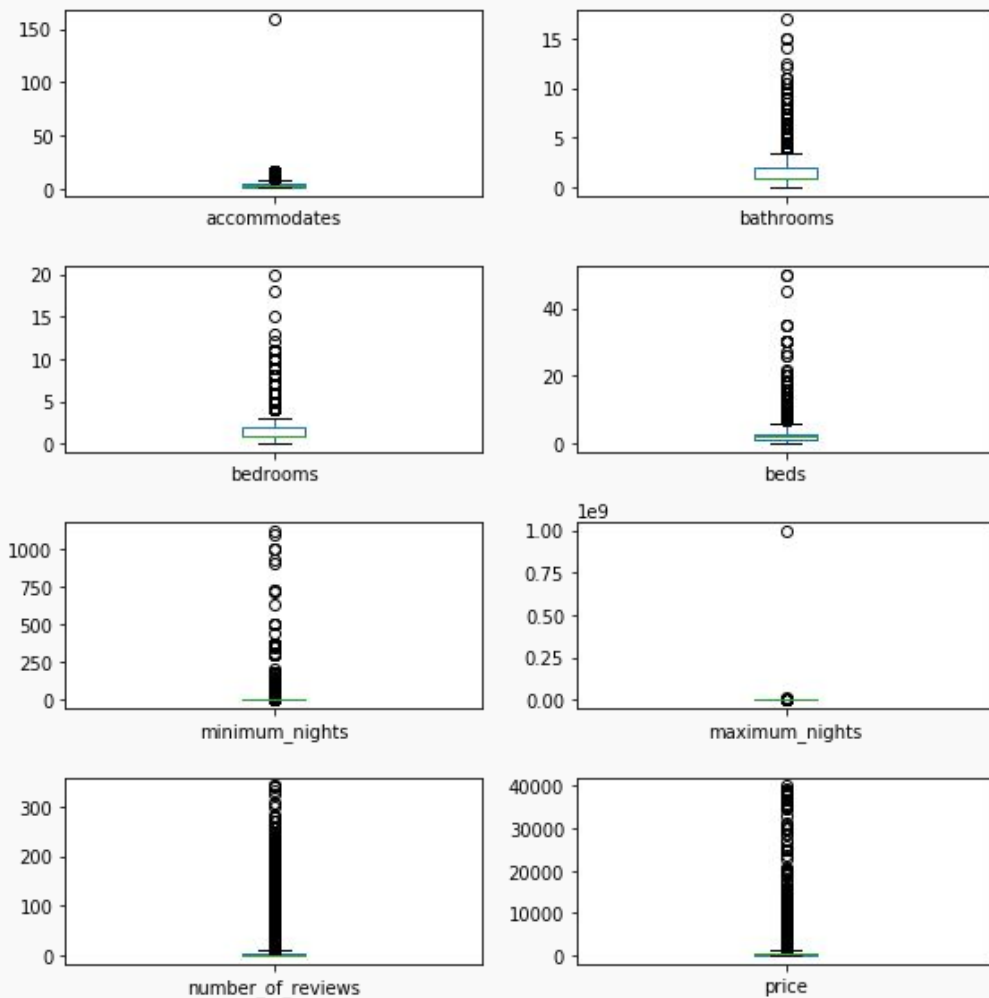

Data Preparation

	accommodates	bathrooms	bedrooms	beds	price	minimum_nights	maximum_nights	number_of_reviews
15364	6	3.0	3.0	3.0	\$1,501.00	1	1125	0
33163	2	2.0	1.0	1.0	\$181.00	5	10	0
28598	3	1.0	1.0	2.0	\$140.00	1	1125	1
24716	1	1.0	1.0	1.0	\$108.00	1	31	0
10299	6	2.0	3.0	5.0	\$189.00	3	1125	53

Exploratory Data Analysis (EDA)

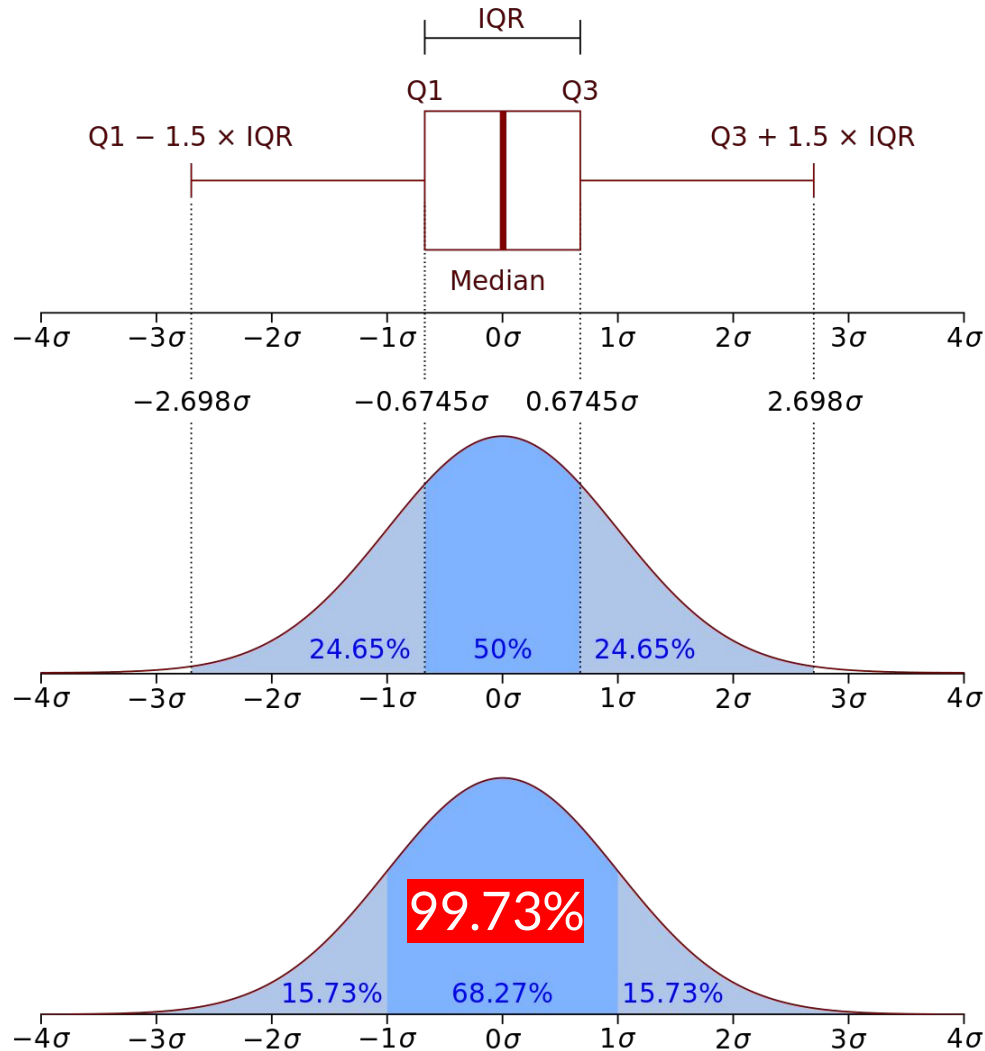


Outliers???

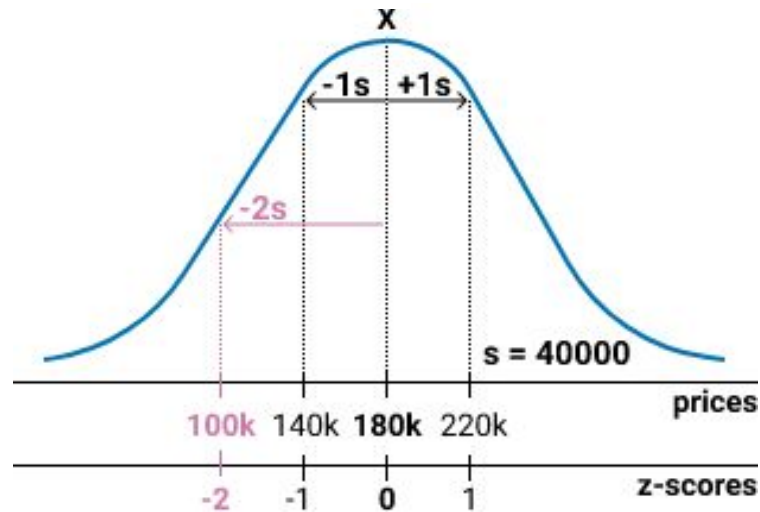


Outliers Removal





Using standardization and IQR for eliminate outlier



$$z = \frac{x - \mu}{\sigma}$$

Standardization (option #01)

	accommodates	bathrooms	bedrooms	beds	minimum_nights	maximum_nights	number_of_reviews	price
0	5	1.0	2.0	2.0	5	30	230	305.0
1	3	1.0	1.0	2.0	4	30	232	158.0
2	3	1.0	1.0	2.0	2	1125	256	251.0
3	3	1.5	1.0	2.0	2	89	155	347.0
4	2	1.0	1.0	1.0	3	28	299	220.0

	accommodates	bathrooms	bedrooms	beds	minimum_nights	maximum_nights	number_of_reviews	price
0	0.309556	-0.665668	0.336640	-0.298117	-0.201519	0.013014	-0.005473	10.148919
1	-0.457286	-0.665668	-0.601309	-0.298117	-0.291600	-0.033019	-0.005473	10.240240
2	-0.457286	-0.665668	-0.601309	-0.298117	-0.234610	-0.125087	-0.005267	11.336088
3	-0.457286	-0.181869	-0.601309	-0.298117	-0.175782	-0.125087	-0.005462	6.724392
4	-0.840707	-0.665668	-0.601309	-0.793011	-0.253607	-0.079053	-0.005474	13.299483

Standardization (mean zero and unitary std)

	accommodates	bathrooms	bedrooms	beds	minimum_nights	maximum_nights	number_of_reviews	price
count	3.563000e+04	3.563000e+04	3.563000e+04	3.563000e+04	3.563000e+04	3.563000e+04	3.563000e+04	3.563000e+04
mean	1.443714e-15	-4.422863e-15	2.209665e-16	7.564037e-16	-8.207401e-16	1.233092e-15	2.345833e-16	9.341216e-15
std	1.000014e+00	1.000014e+00	1.000014e+00	1.000014e+00	1.000014e+00	1.000014e+00	1.000014e+00	1.000014e+00
min	-1.224128e+00	-1.633266e+00	-1.539257e+00	-1.287905e+00	-3.884207e-01	-1.711206e-01	-5.478829e-03	-3.529624e-01
25%	-8.407069e-01	-6.656684e-01	-6.013086e-01	-7.930111e-01	-2.958891e-01	-1.711206e-01	-5.473355e-03	-3.529624e-01
50%	-7.386478e-02	-6.656684e-01	-6.013086e-01	-2.981171e-01	-2.113239e-01	-1.250869e-01	-5.266671e-03	-3.073020e-01
75%	3.095563e-01	3.019297e-01	3.366401e-01	1.967770e-01	-2.197118e-02	-3.301947e-02	-5.266671e-03	-1.703209e-01
max	5.973982e+01	1.481590e+01	1.721972e+01	2.345680e+01	2.412264e+01	5.147871e+01	1.887469e+02	1.521722e+01

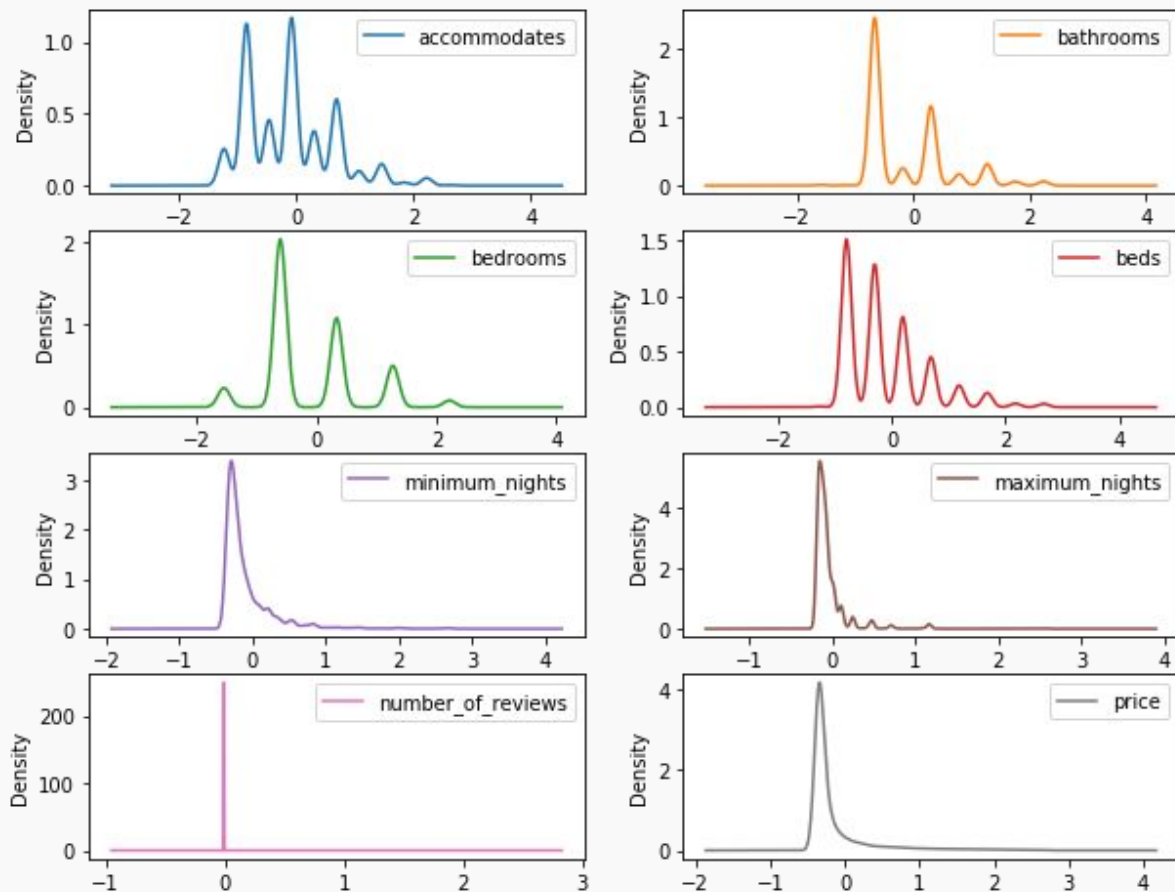
```
from sklearn.preprocessing import StandardScaler

# get data
target_columns = ["accommodates", "bathrooms", "bedrooms",
                  "beds", "minimum_nights",
                  "maximum_nights", "number_of_reviews", "price"]
rio_listings = pd.read_csv("listings.csv", usecols=target_columns)



# clean missing values
rio_listings.dropna(axis=0, inplace=True)

# apply z-score (mean=0, std=1)
rio_z_scored = pd.DataFrame(StandardScaler().fit_transform(rio_listings),
                             columns=target_columns,
                             index=rio_listings.index)

# remove outliers
rio_z_scored = rio_z_scored[(rio_z_scored < 2.698).all(axis=1)
                             & (rio_z_scored > -2.698).all(axis=1)]
```

After
Standardization



```
Q1 = rio_iqr.quantile(0.25)
```

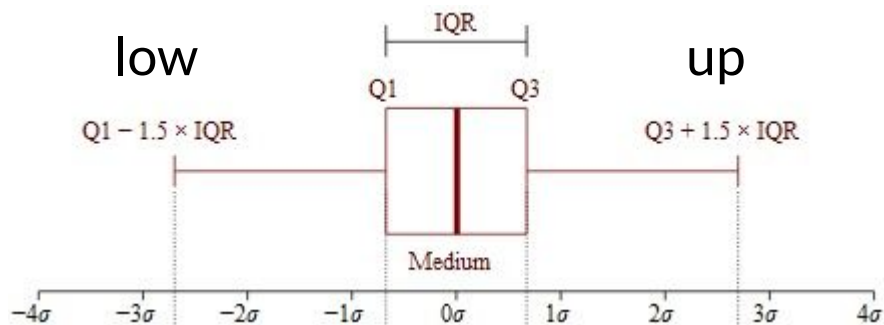
accommodates	2.0
bathrooms	1.0
bedrooms	1.0
beds	1.0
minimum_nights	1.0
maximum_nights	30.0
number_of_reviews	0.0
price	151.0

```
IQR = Q3 - Q1
```

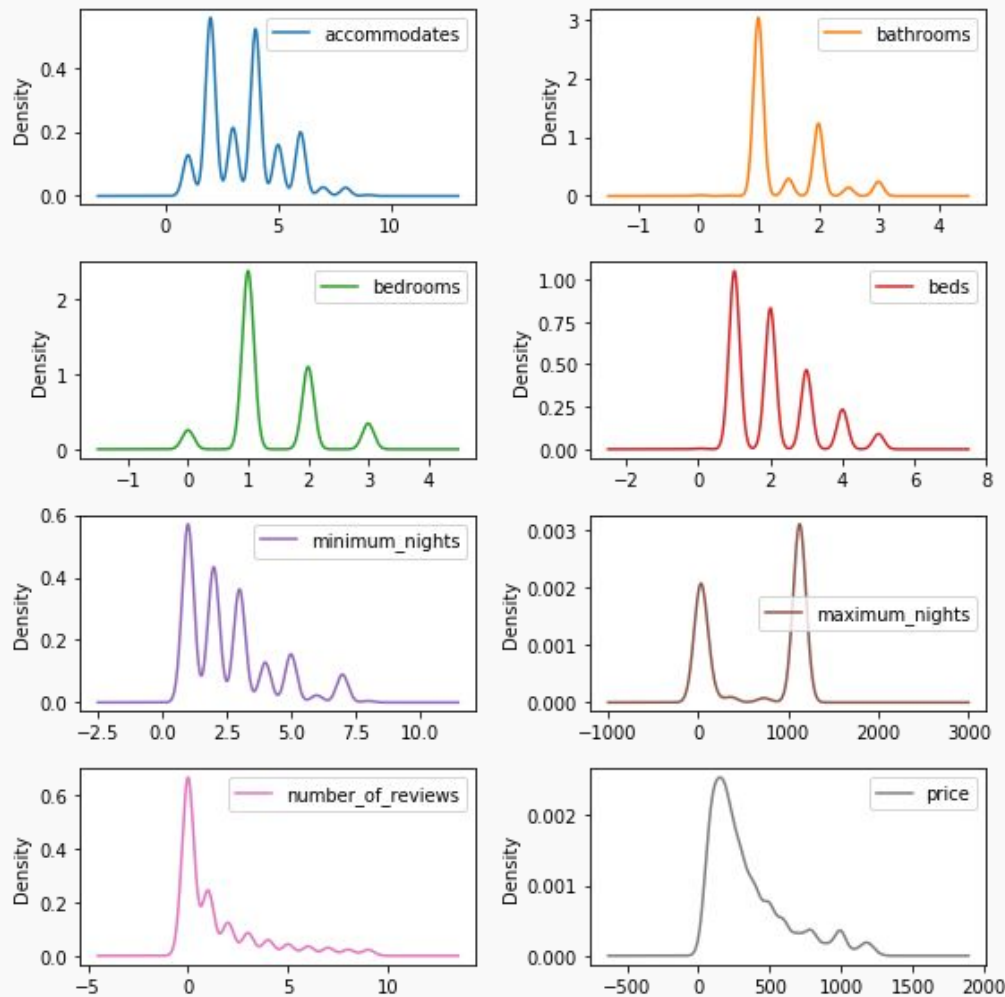
accommodates	3.0
bathrooms	1.0
bedrooms	1.0
beds	2.0
minimum_nights	3.0
maximum_nights	1095.0
number_of_reviews	4.0
price	447.0

```
Q3 = rio_iqr.quantile(0.75)
```

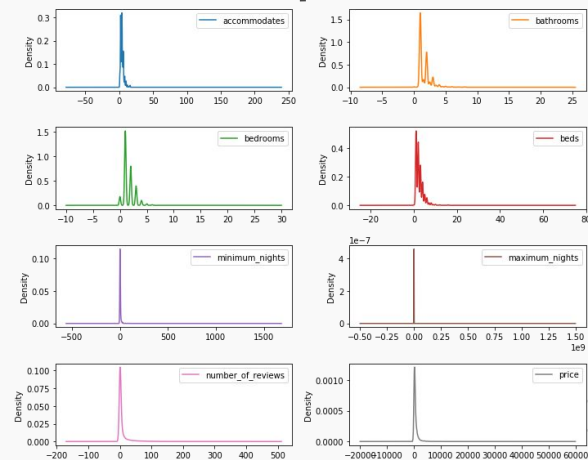
accommodates	5.0
bathrooms	2.0
bedrooms	2.0
beds	3.0
minimum_nights	4.0
maximum_nights	1125.0
number_of_reviews	4.0
price	598.0



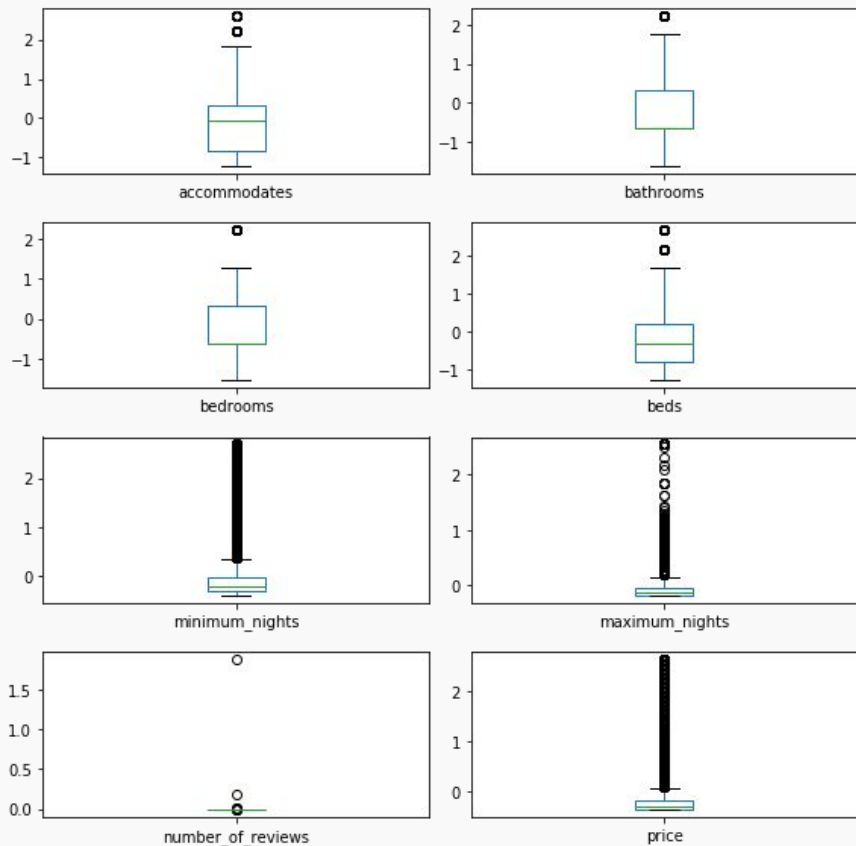
```
rio_iqr = rio_listings[target_columns].copy()
```



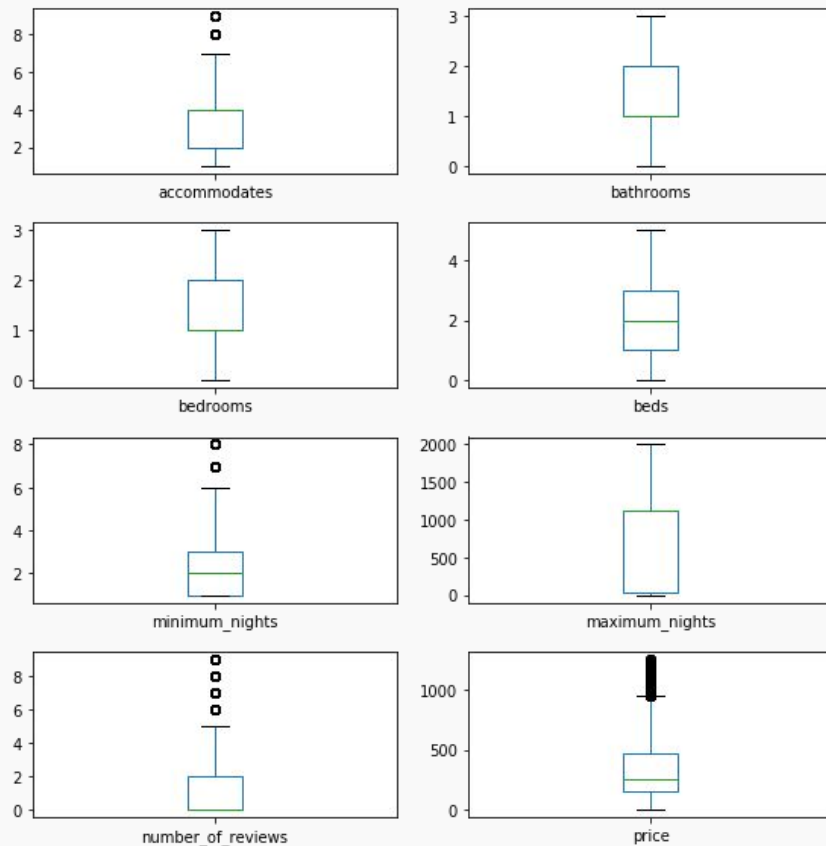
After outlier
elimination (IQR
method)



Z-Score Method




IQR Method

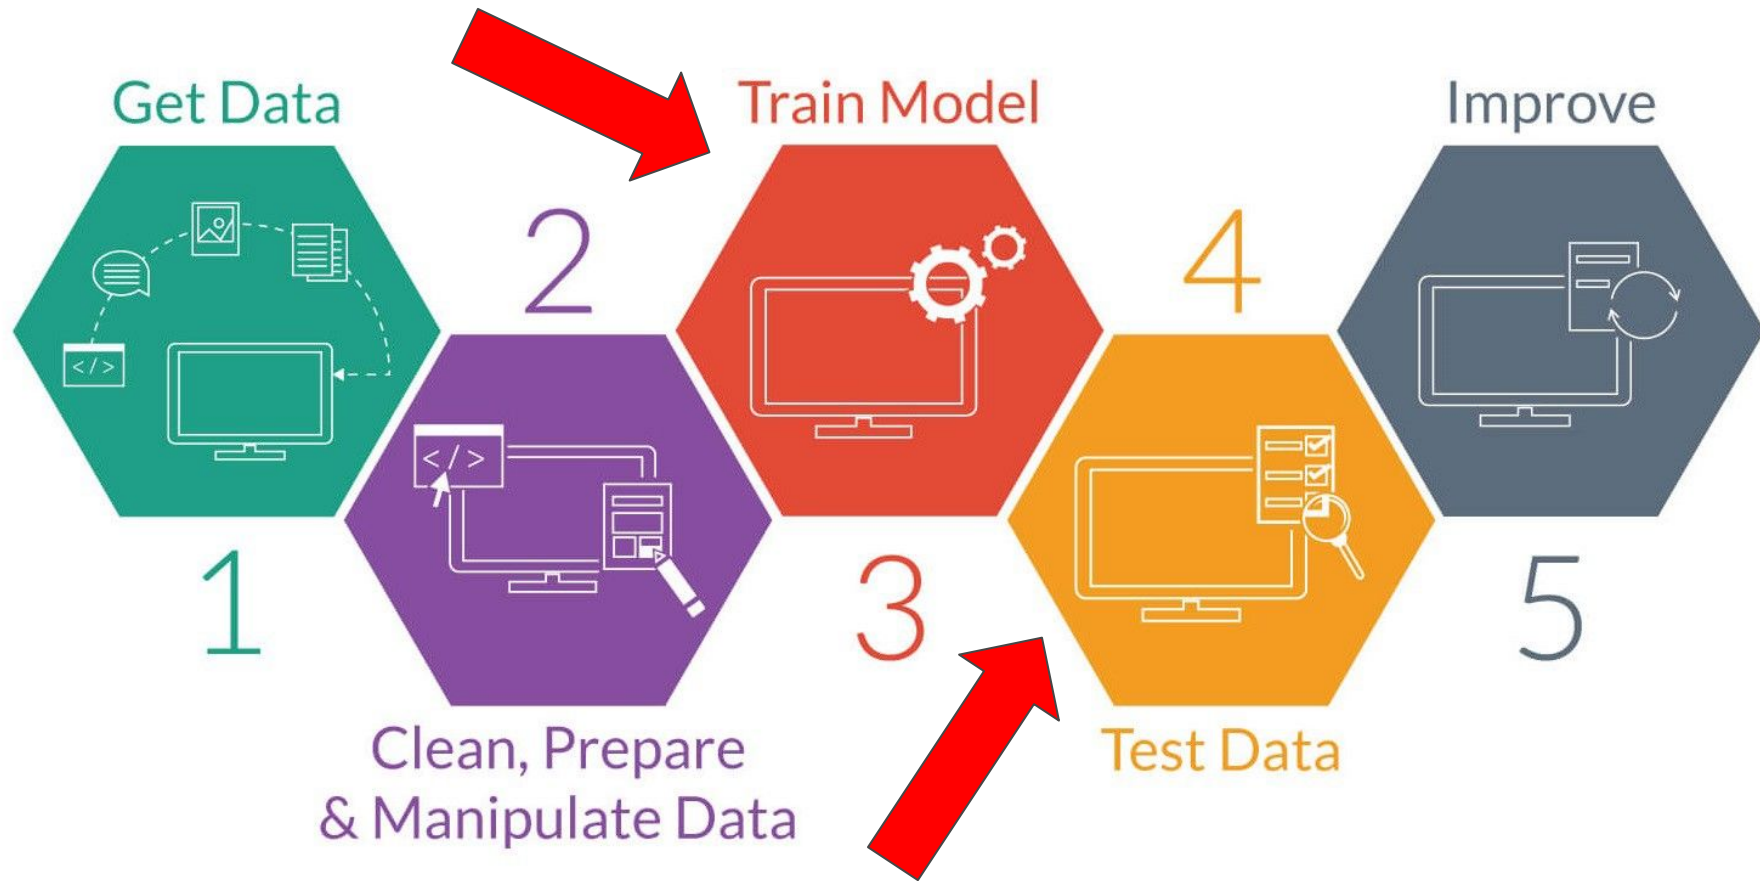


Separate data into Training and Test


```
X_train, X_test, Y_train, Y_test = train_test_split(rio_iqr,  
                                                    rio_iqr["price"],  
                                                    test_size=test_size,  
                                                    random_state=seed)
```



A general ML workflow



Introduction to scikit-learn



Classification

Identifying to which category an object belongs to.

Applications: Spam detection, Image recognition.

Algorithms: SVM, nearest neighbors, random forest, ... — Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: SVR, ridge regression, Lasso, ... — Examples

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, ... — Examples

Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency

Algorithms: PCA, feature selection, non-negative matrix factorization. — Examples

Model selection

Comparing, validating and choosing parameters and models.

Goal: Improved accuracy via parameter tuning

Modules: grid search, cross validation, metrics. — Examples

Preprocessing

Feature extraction and normalization.

Application: Transforming input data such as text for use with machine learning algorithms.

Modules: preprocessing, feature extraction. — Examples

Scikit-learn workflow

The scikit-learn workflow consists of 4 main steps:

- instantiate the specific machine learning model you want to use
- fit the model to the training data
- use the model to make predictions
- evaluate the accuracy of the predictions

```
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error
import numpy as np
```

```
# instantiate a knn object
```

```
knn = KNeighborsRegressor(n_neighbors=4, n_jobs=-1)
```

```
# train the model
```

```
knn.fit(X_train,Y_train)
```

```
# predict
```


```
predict = knn.predict(X_test)
```

```
# evaluate
```

```
rmse = np.sqrt(mean_squared_error(Y_test,predict))
```

2.5381

K-Fold Cross Validation




```
# Test options and evaluation metric
```

```
num_folds = 10
```

```
seed = 7
```

```
scoring = 'neg_mean_squared_error'
```

```
# Standardize the dataset
```

```
pipelines = []
```

```
pipelines.append(('NonScaledKnn',  
                  Pipeline([('KNN',  
                             KNeighborsRegressor(n_neighbors=5, n_jobs=-1))])))
```

```
pipelines.append(('ScaledKnn',  
                  Pipeline([('Scaler',  
                             StandardScaler()),  
                             ('KNN',  
                              KNeighborsRegressor(n_neighbors=5, n_jobs=-1))])))
```

```
pipelines.append(('NormalizedKnn',  
                  Pipeline([('Normalizer',  
                             Normalizer()),  
                             ('KNN',  
                              KNeighborsRegressor(n_neighbors=5, n_jobs=-1))])))
```

```
results = []
names = []
for name, model in pipelines:
    kfold = KFold(n_splits=num_folds, random_state=seed)
    cv_results = cross_val_score(model, X_train, Y_train, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    print("%s Mean: %f Std: %f" % (name,
                                   np.sqrt(-cv_results.mean()),
                                   np.sqrt(cv_results.std())))
```

NonScaledKnn Mean: 2.532367 Std: 1.713266

ScaledKnn Mean: 48.808902 Std: 9.812108


NormalizedKnn Mean: 93.813966 Std: 26.146305

RobustedKnn Mean: 42.450476 Std: 7.967048

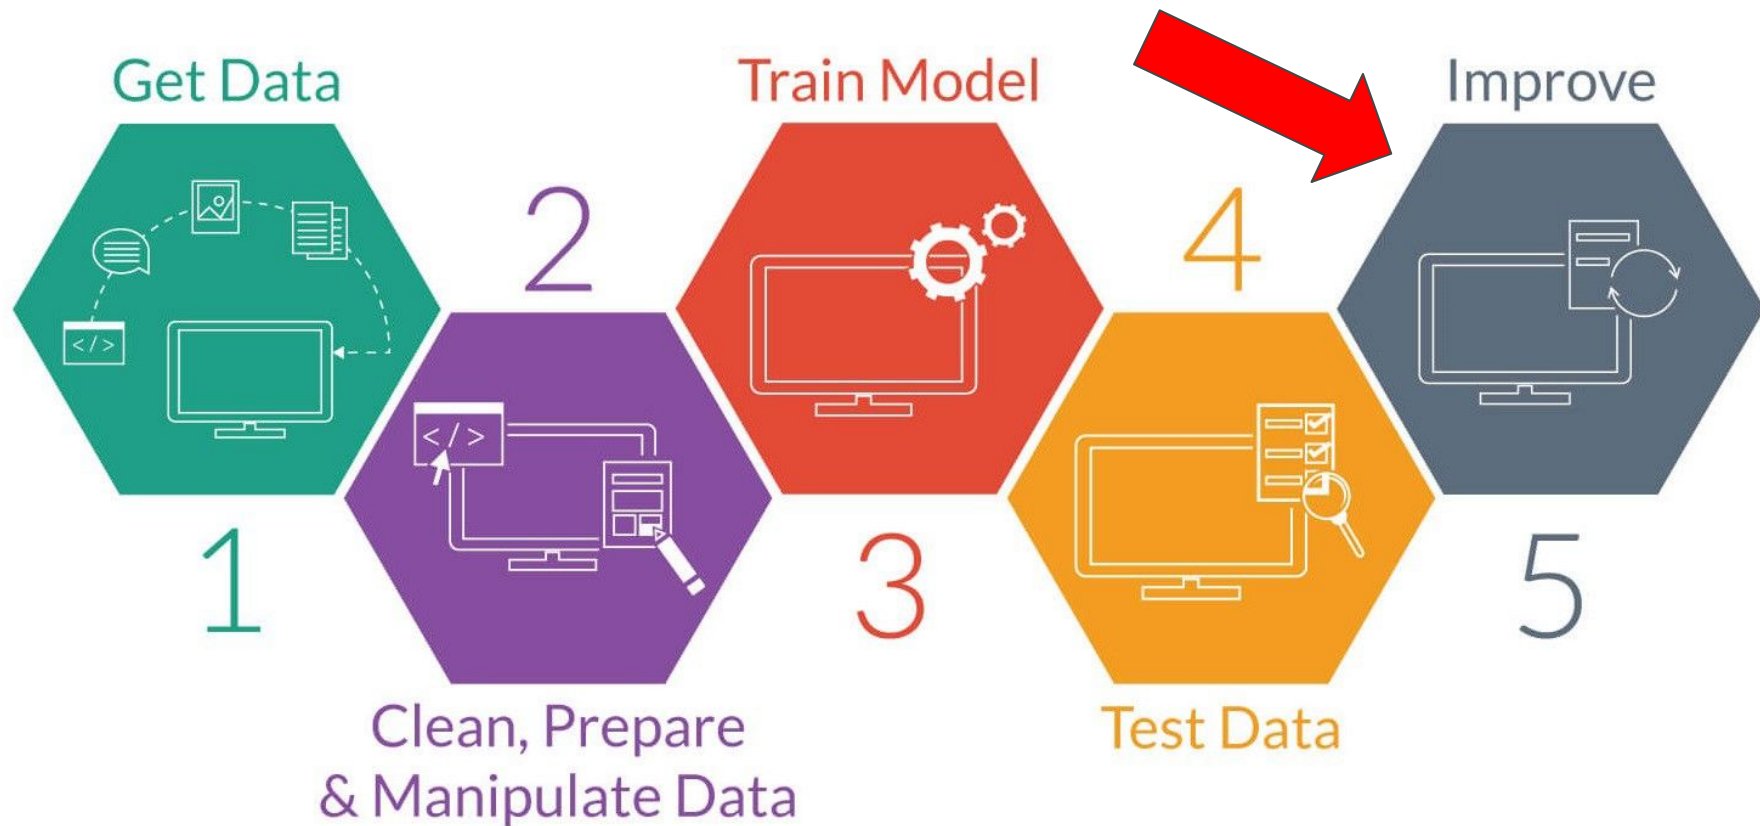
QuantiledKnn Mean: 72.075350 Std: 20.764550

PoweredKnn Mean: 58.408967 Std: 12.822987

Scaled Algorithm Comparison



A general ML workflow



Improve the accuracy

- **Increase the number of attributes** the model uses to calculate similarity when ranking the closest neighbors
- When we **vary the features** that are used in the model, we're **affecting the data** that the model uses.
- **Increase k** , the number of nearby neighbors the model uses when computing the prediction
- On the other hand, **varying the k value affects the behavior of the model independently of the actual data** that's used when making predictions. Values that affect the behavior and performance of a model that are unrelated to the data that's used are referred to as **hyperparameters**.

Hyperparameter Optimization

A simple but common [hyperparameter optimization](#) technique is known as [grid search](#):

- selecting a subset of the possible hyperparameter values,
- training a model using each of these hyperparameter values,
- evaluating each model's performance,
- selecting the hyperparameter value that resulted in the lowest error value.


```
# hyperparameter
k_values = np.array([1,3,5,7,9,11,13,15,17,19,21])
param_grid = dict(n_neighbors=k_values)

# scaler
scaler = RobustScaler().fit_transform(X_train)

# instantiate a model
model = KNeighborsRegressor()

# Test options and evaluation metric
num_folds = 10
seed = 20
scoring = 'neg_mean_squared_error'

# cross-validation
kfold = KFold(n_splits=num_folds, random_state=seed)
grid = GridSearchCV(estimator=model,
                    param_grid=param_grid,
                    scoring=scoring,
                    cv=kfold)

grid_result = grid.fit(scaler, Y_train)
```

```
Best: 42.266050 using {'n_neighbors': 7}
57.668295 (10.088662) with: {'n_neighbors': 1}
45.151191 (7.902694) with: {'n_neighbors': 3}
42.606809 (8.416300) with: {'n_neighbors': 5}
42.266050 (9.352537) with: {'n_neighbors': 7}
42.572611 (9.835827) with: {'n_neighbors': 9}
43.060314 (9.638778) with: {'n_neighbors': 11}
43.604723 (9.396493) with: {'n_neighbors': 13}
44.049894 (9.454142) with: {'n_neighbors': 15}
44.498295 (9.868717) with: {'n_neighbors': 17}
44.944889 (9.781728) with: {'n_neighbors': 19}
45.591393 (10.107396) with: {'n_neighbors': 21}
```


Finalize Model

A white puzzle piece is missing from a white puzzle, revealing a red surface underneath. The puzzle piece is located in the lower right quadrant of the image.

```
predict = grid_result.best_estimator_.predict(RobustScaler().fit_transform(X_test))  
rmse = np.sqrt(mean_squared_error(predict,Y_test))
```

49.44294556465665

Bias-Variance Tradeoff





```
index.js
import React, { useState } from 'react';
import ReactDOM from 'react-dom';

const App = () => {
  const [contacts, setContacts] = useState([]);

  const addContact = (e) => {
    e.preventDefault();
    const name = document.getElementById('name').value;
    const phone = document.getElementById('phone').value;
    setContacts([...contacts, { name, phone }]);
  };

  return (
    <div>
      <input type="text" value="" id="name" />
      <input type="text" value="" id="phone" />
      <button onClick={addContact}>Add Contact</button>
    </div>
  );
};

ReactDOM.render(<App />, document.getElementById('root'));
```

```
index.html
<div>
  <script src="index.js"></script>
  <script src="https://unpkg.com/react@17.0.2/umd/react.development.js"></script>
  <script src="https://unpkg.com/react-dom@17.0.2/umd/react-dom.development.js"></script>
  <script src="https://unpkg.com/babel-standalone@6.26.0/babel.js"></script>
  <script src="https://unpkg.com/react-scripts@1.4.2/dist/index.js"></script>
  <div id="root"></div>
</div>
```