

Lesson #04 - Linear Regression

- We are going to start by covering **linear regression**
- We discuss the application of linear regression to **housing price prediction**
- Present the notion of a **cost function**
- Introduce the **gradient descent** method for learning.
- Refresher on **linear algebra concepts**.

Instance-Based
Learning
Vs
Model-Based
Learning

Instance-Based Learning



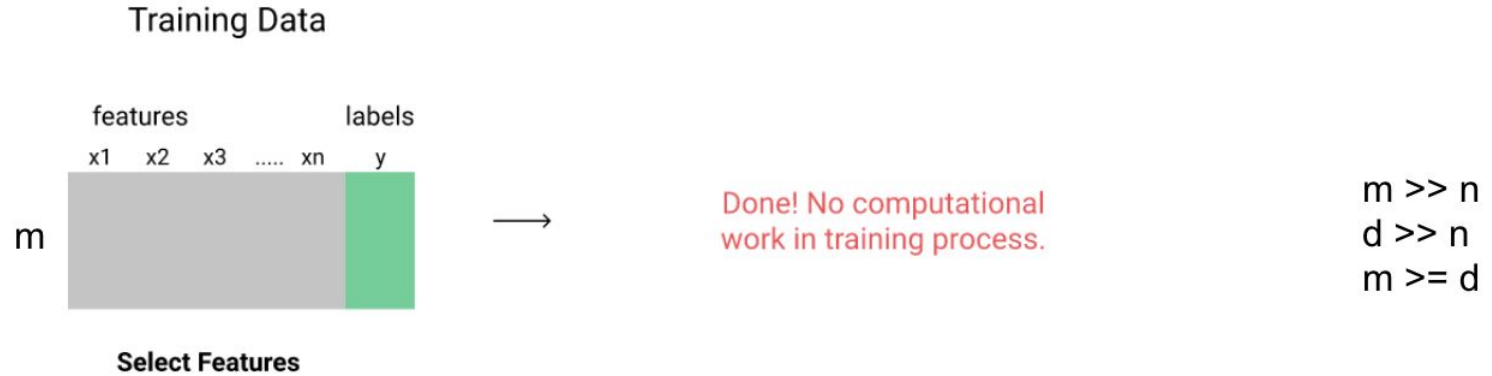
Model-Based Learning

$$f(x, \alpha, \beta)$$

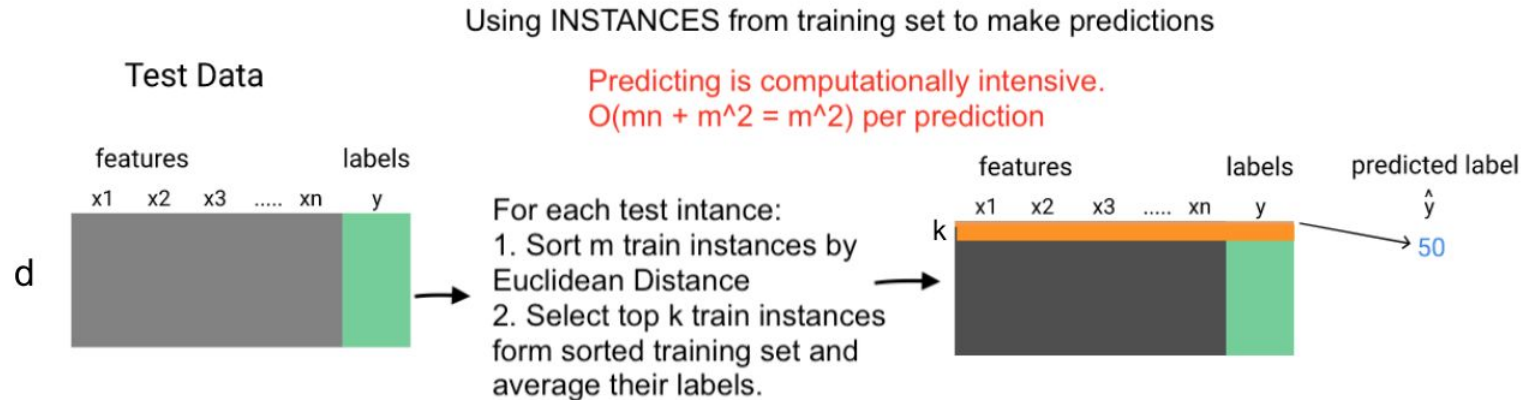
Instance-Based Learning: KNN

4

Training Process



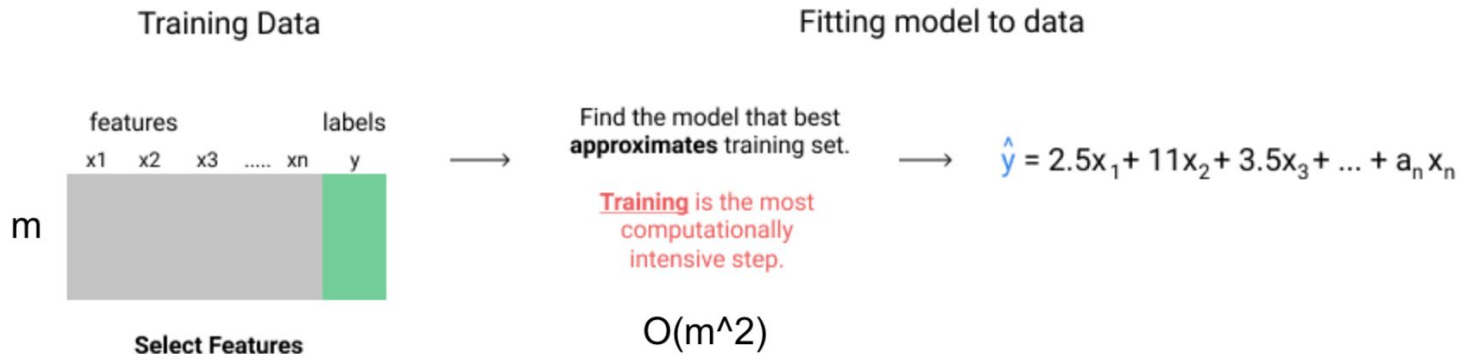
Testing Process



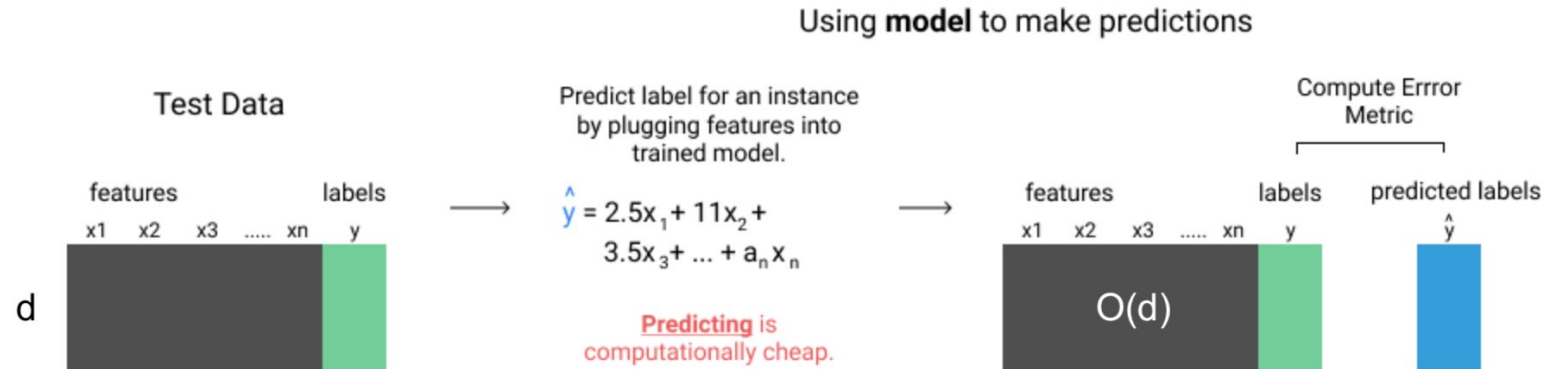
Model-Based Learning: Linear Regression

5

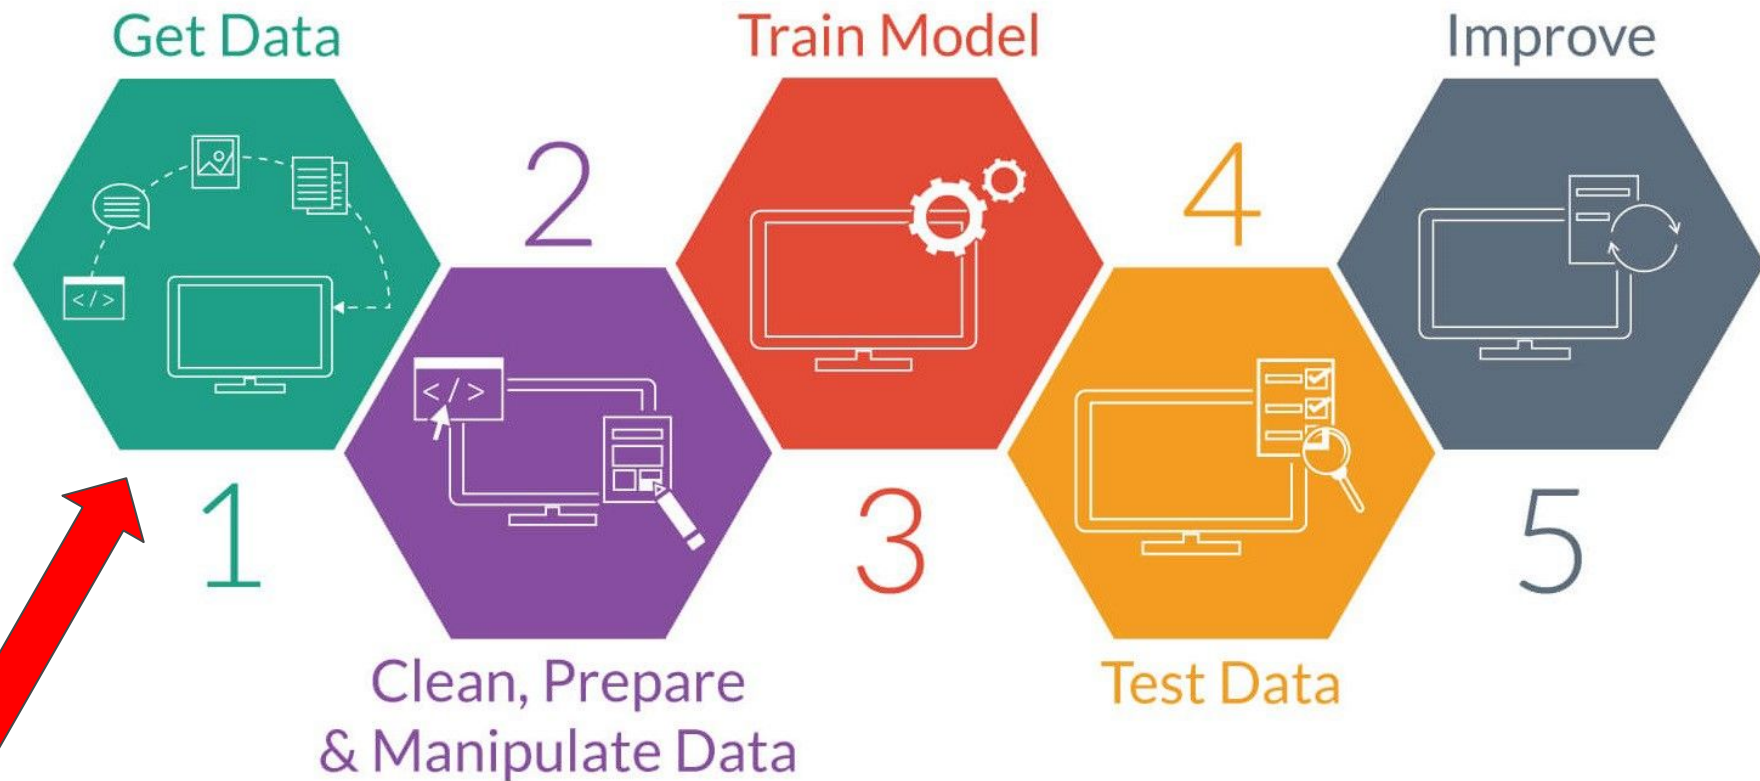
Training Process



Testing Process



A general ML workflow





Ames, Iowa: Alternative to the Boston Housing Data as an End of Semester Regression Project

[Dean De Cock](#)

Truman State University

Journal of Statistics Education Volume 19, Number 3(2011),
www.amstat.org/publications/jse/v19n3/decock.pdf

Copyright © 2011 by Dean De Cock all rights reserved. This text may be freely shared among individuals, but it may not be republished in any medium without express written consent from the author and advance notification of the editor.

Key Words: Multiple Regression; Linear Models; Assessed Value; Group Project.

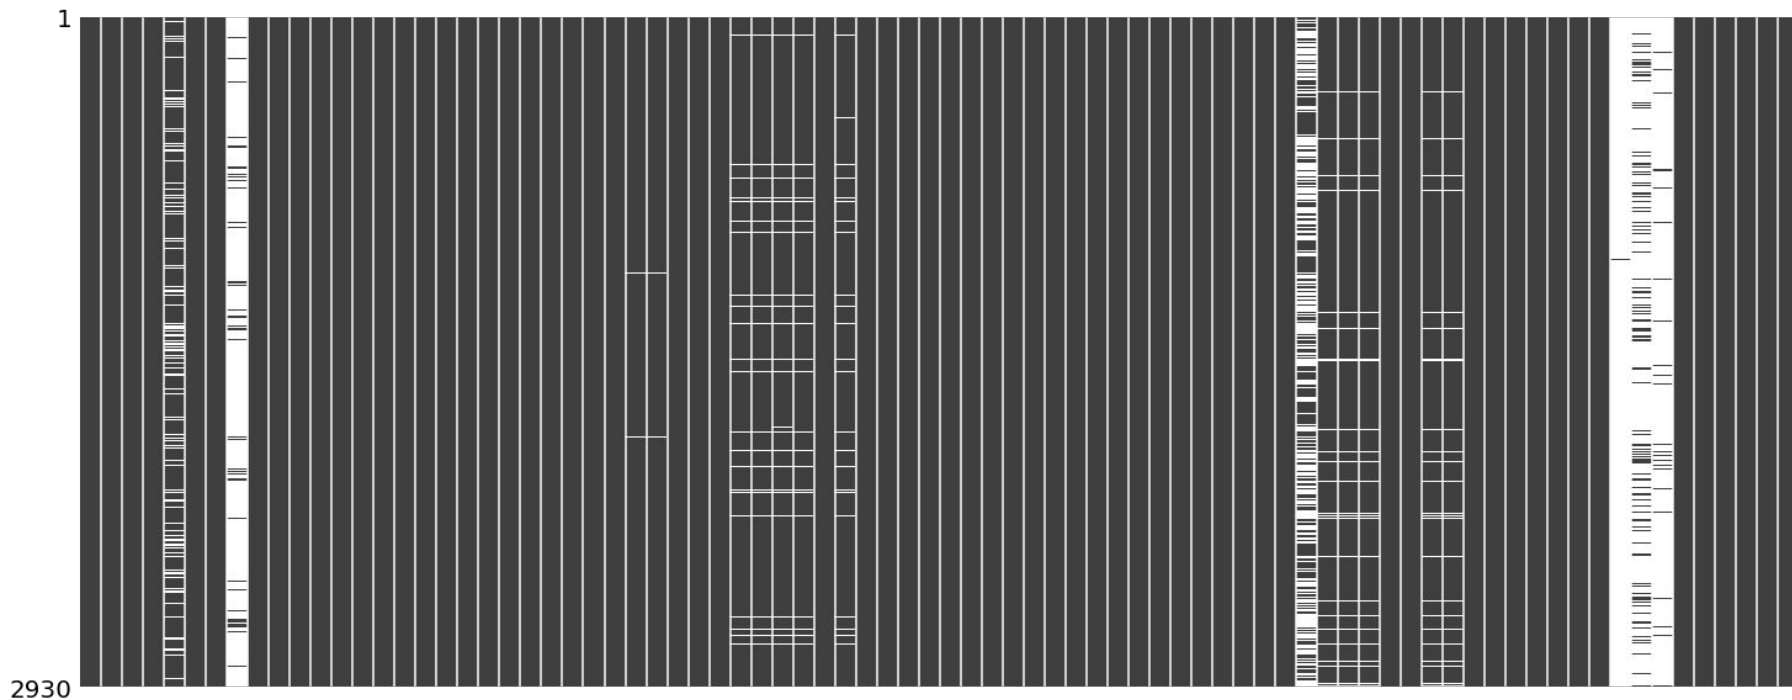
Abstract

This paper presents a data set describing the sale of individual residential property in Ames, Iowa from 2006 to 2010. The data set contains 2930 observations and a large number of explanatory variables (23 nominal, 23 ordinal, 14 discrete, and 20 continuous) involved in assessing home values. I will discuss my previous use of the Boston Housing Data Set and I will suggest methods for incorporating this new data set as a final project in an undergraduate regression course.

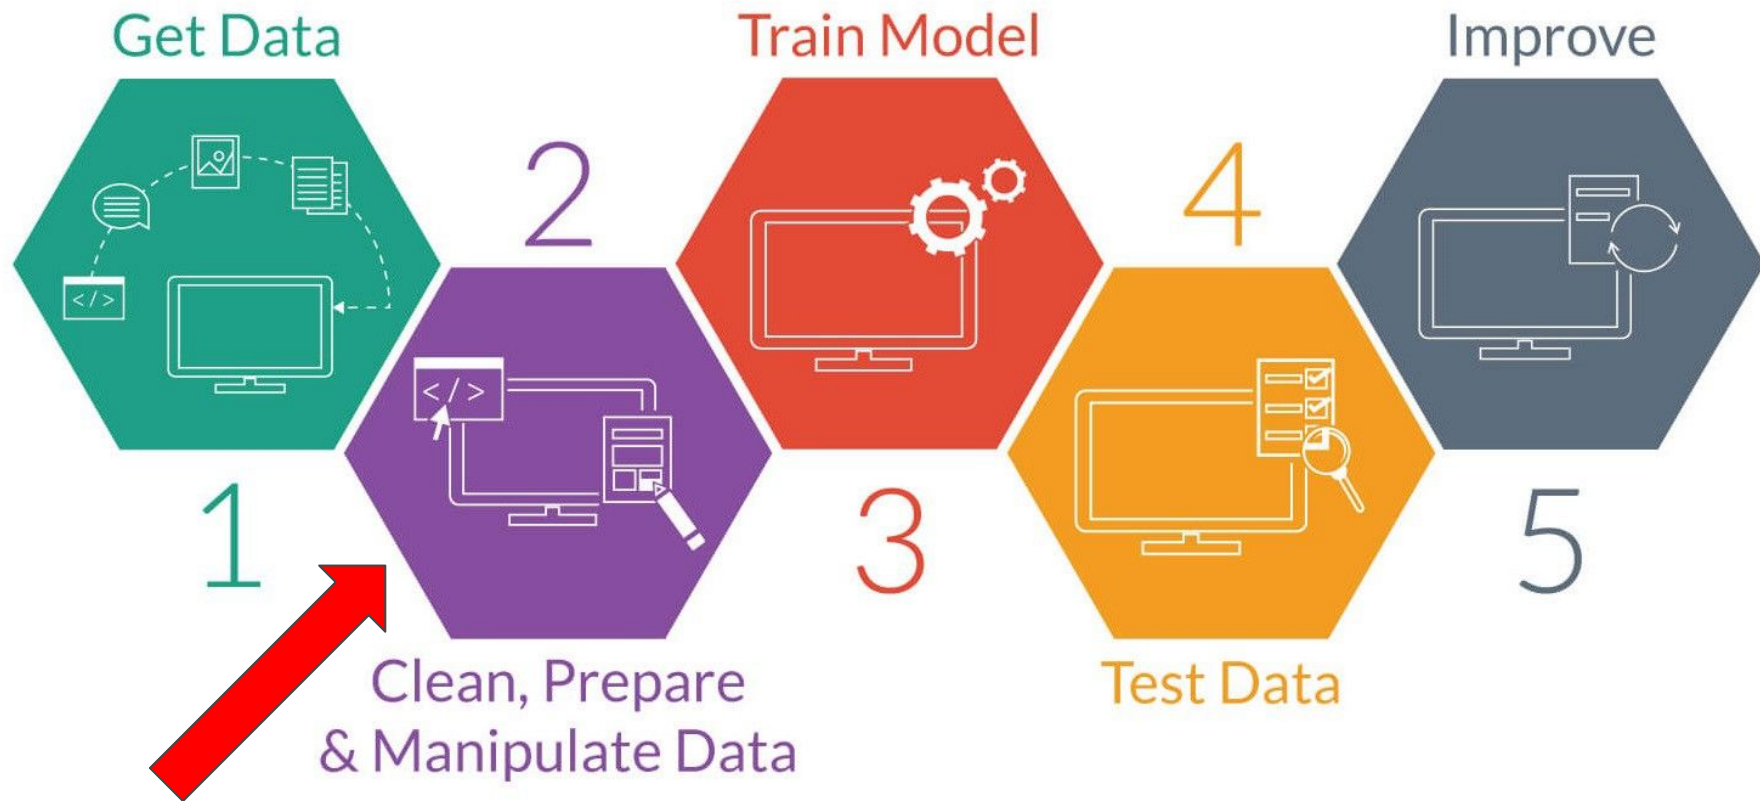
(2930, 82)

Visualizing Missing Values

```
import pandas as pd
import missingno as msno
# get the data
data = pd.read_csv("AmesHousing.txt", sep='\t')
# visualize missing values
msno.matrix(data, sparkline=False)
```



A general ML workflow

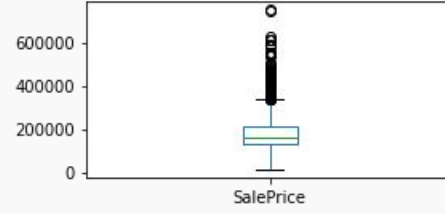
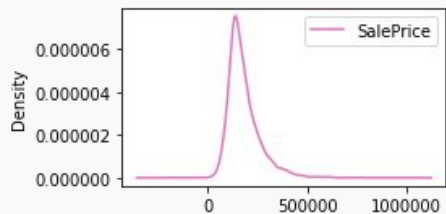
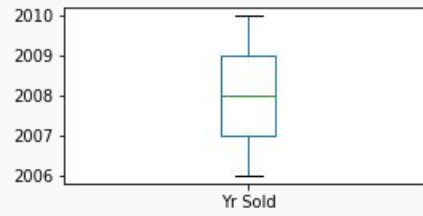
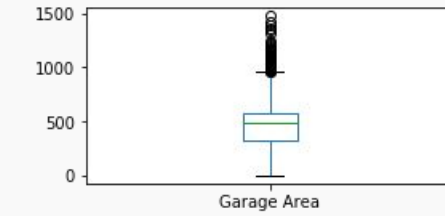
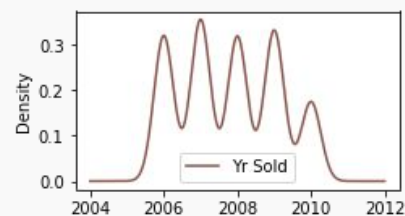
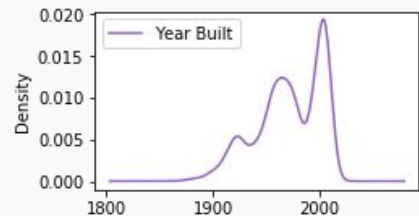
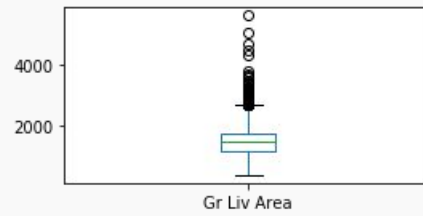
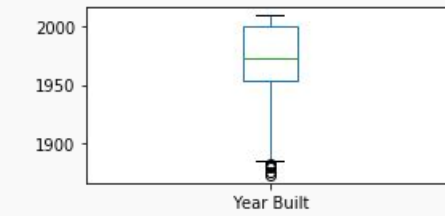
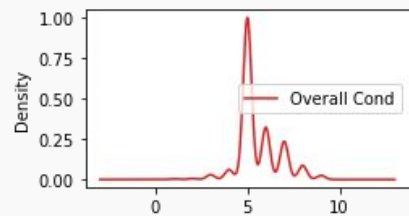
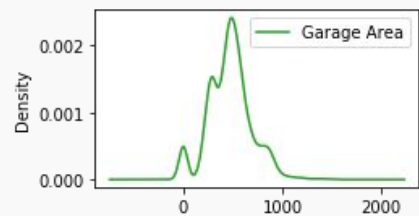
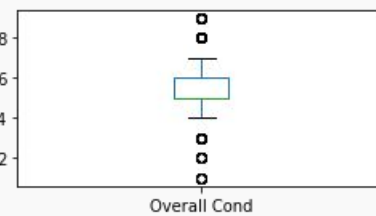
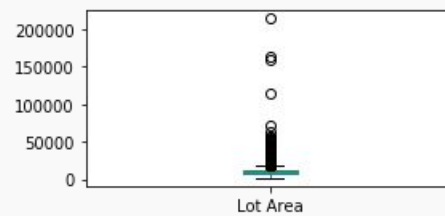
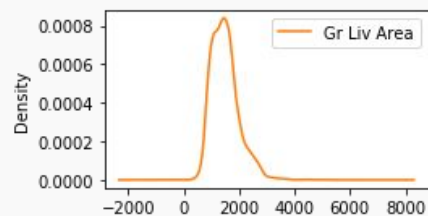
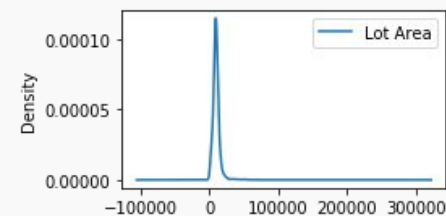


| | Lot Area | Gr Liv Area | Garage Area | Overall Cond | Year Built | Yr Sold | SalePrice |
|------|----------|-------------|-------------|--------------|------------|---------|-----------|
| 1118 | 9428 | 1874 | 880.0 | 5 | 2007 | 2008 | 297900 |
| 2786 | 9800 | 894 | 552.0 | 7 | 1972 | 2006 | 149900 |
| 2633 | 7000 | 864 | 336.0 | 6 | 1962 | 2006 | 105000 |
| 2002 | 9439 | 1248 | 160.0 | 5 | 1930 | 2007 | 87000 |
| 268 | 4435 | 848 | 420.0 | 5 | 2003 | 2010 | 143750 |

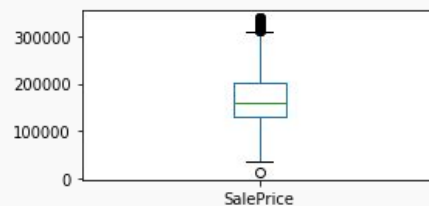
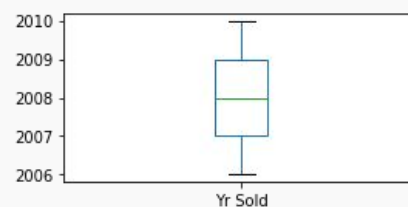
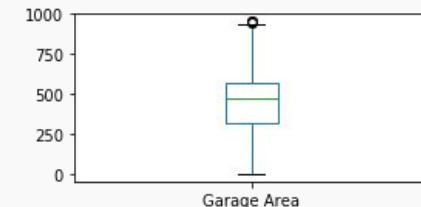
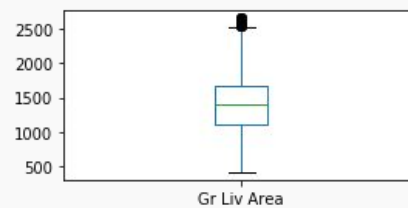
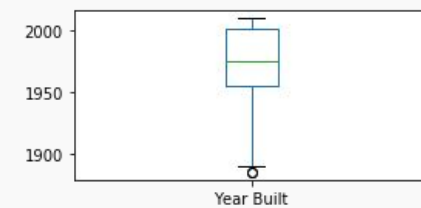
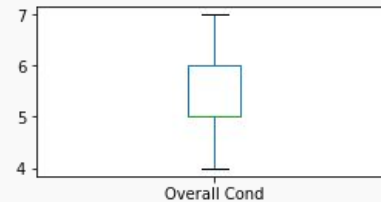
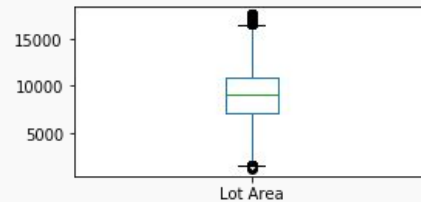
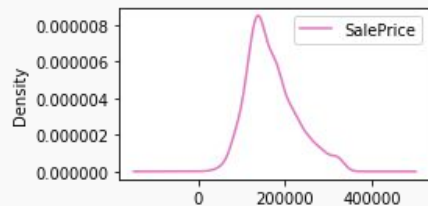
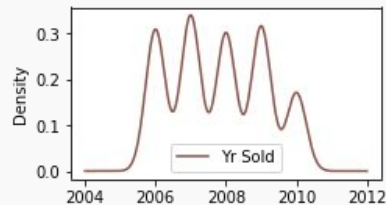
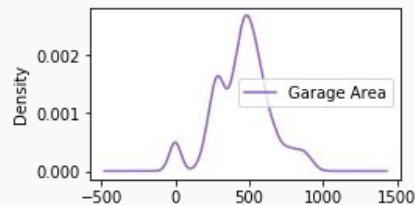
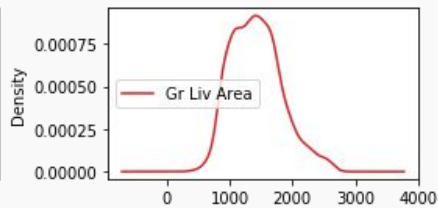
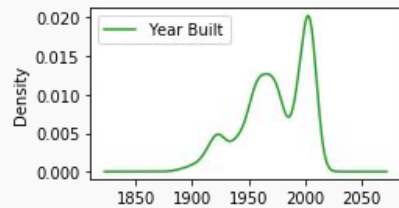
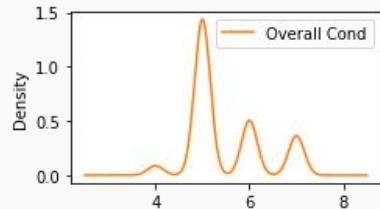
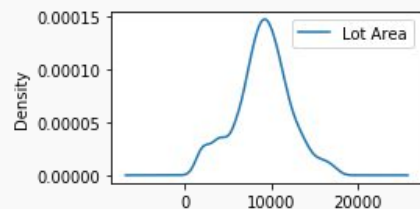
EDA

| | Lot Area | Gr Liv Area | Garage Area | Overall Cond | Year Built | Yr Sold | SalePrice |
|-------|---------------|-------------|-------------|--------------|-------------|-------------|---------------|
| count | 2930.000000 | 2930.000000 | 2929.000000 | 2930.000000 | 2930.000000 | 2930.000000 | 2930.000000 |
| mean | 10147.921843 | 1499.690444 | 472.819734 | 5.563140 | 1971.356314 | 2007.790444 | 180796.060068 |
| std | 7880.017759 | 505.508887 | 215.046549 | 1.111537 | 30.245361 | 1.316613 | 79886.692357 |
| min | 1300.000000 | 334.000000 | 0.000000 | 1.000000 | 1872.000000 | 2006.000000 | 12789.000000 |
| 25% | 7440.250000 | 1126.000000 | 320.000000 | 5.000000 | 1954.000000 | 2007.000000 | 129500.000000 |
| 50% | 9436.500000 | 1442.000000 | 480.000000 | 5.000000 | 1973.000000 | 2008.000000 | 160000.000000 |
| 75% | 11555.250000 | 1742.750000 | 576.000000 | 6.000000 | 2001.000000 | 2009.000000 | 213500.000000 |
| max | 215245.000000 | 5642.000000 | 1488.000000 | 9.000000 | 2010.000000 | 2010.000000 | 755000.000000 |

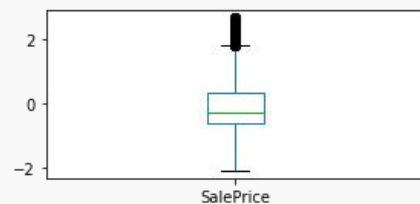
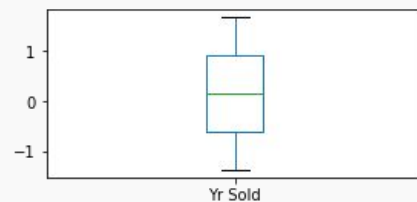
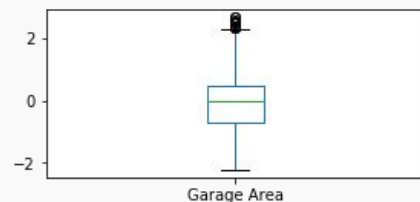
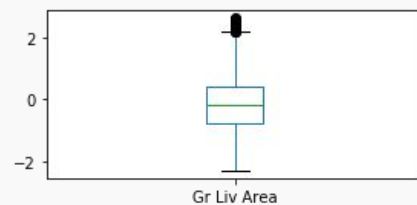
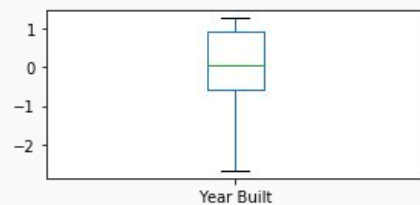
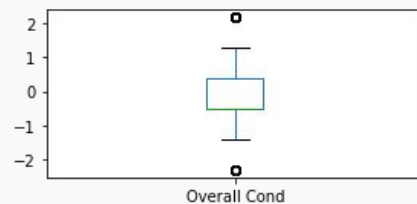
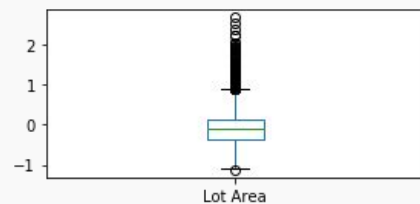
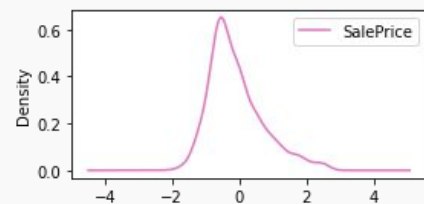
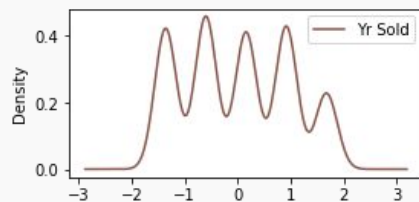
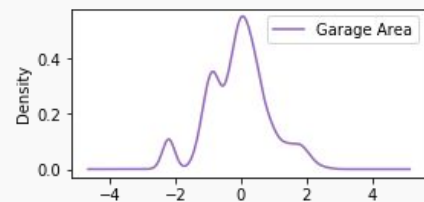
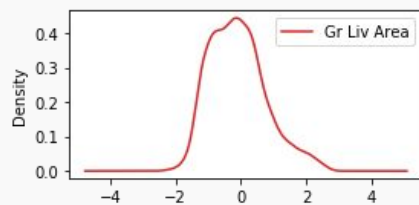
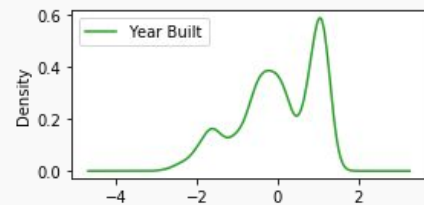
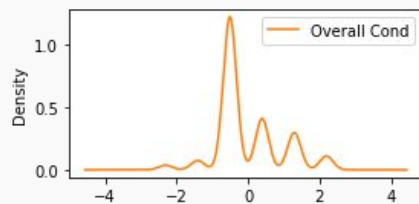
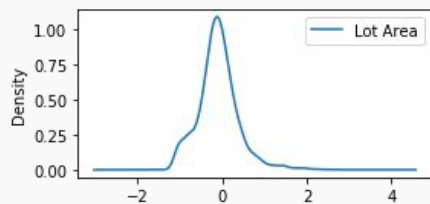
Original Data (2929,7)



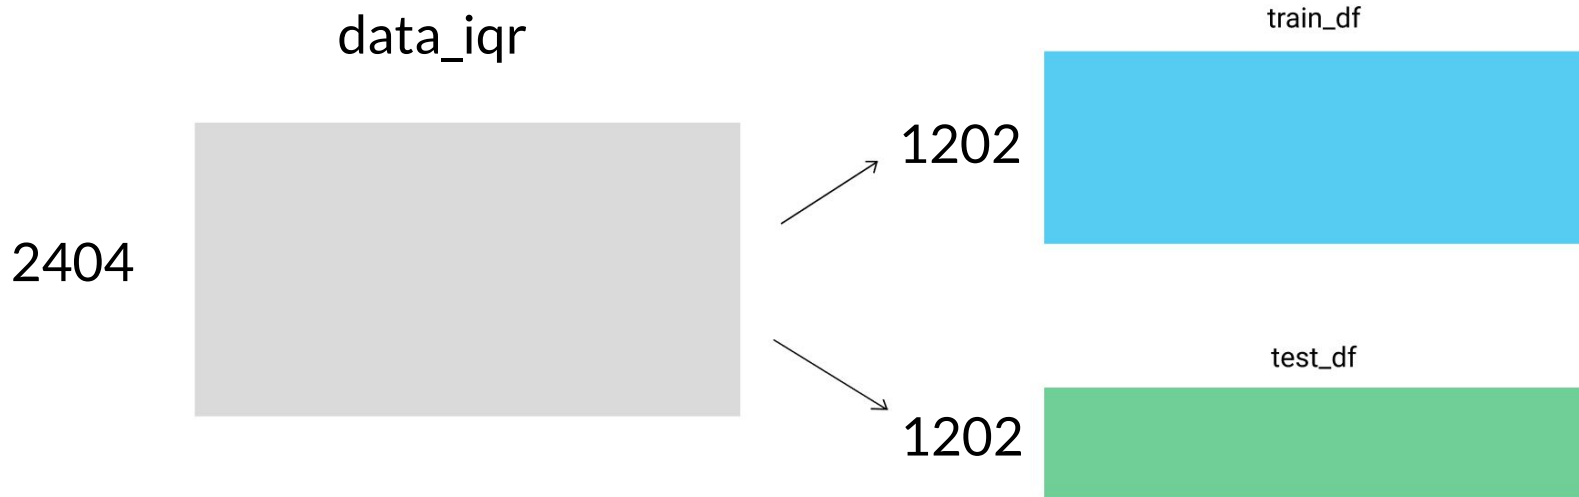
Outlier elimination - IQR method (2404,7) 17.90% of original data were removed



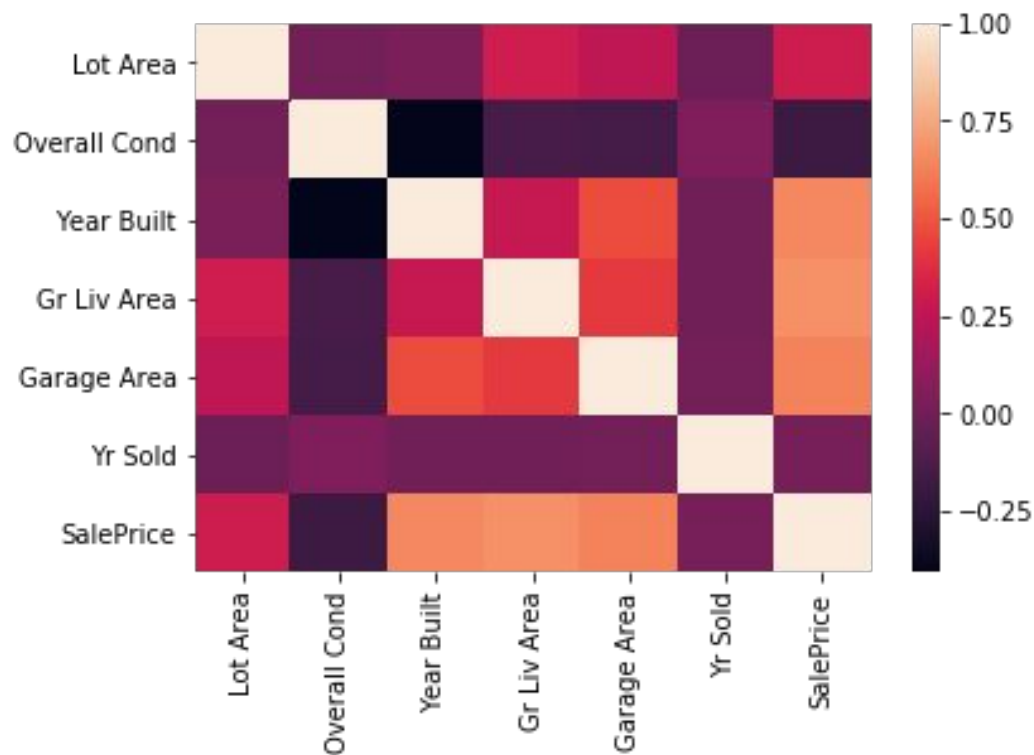
Outlier elimination - Z-Score method (2745,7) 6.28% of original data were removed



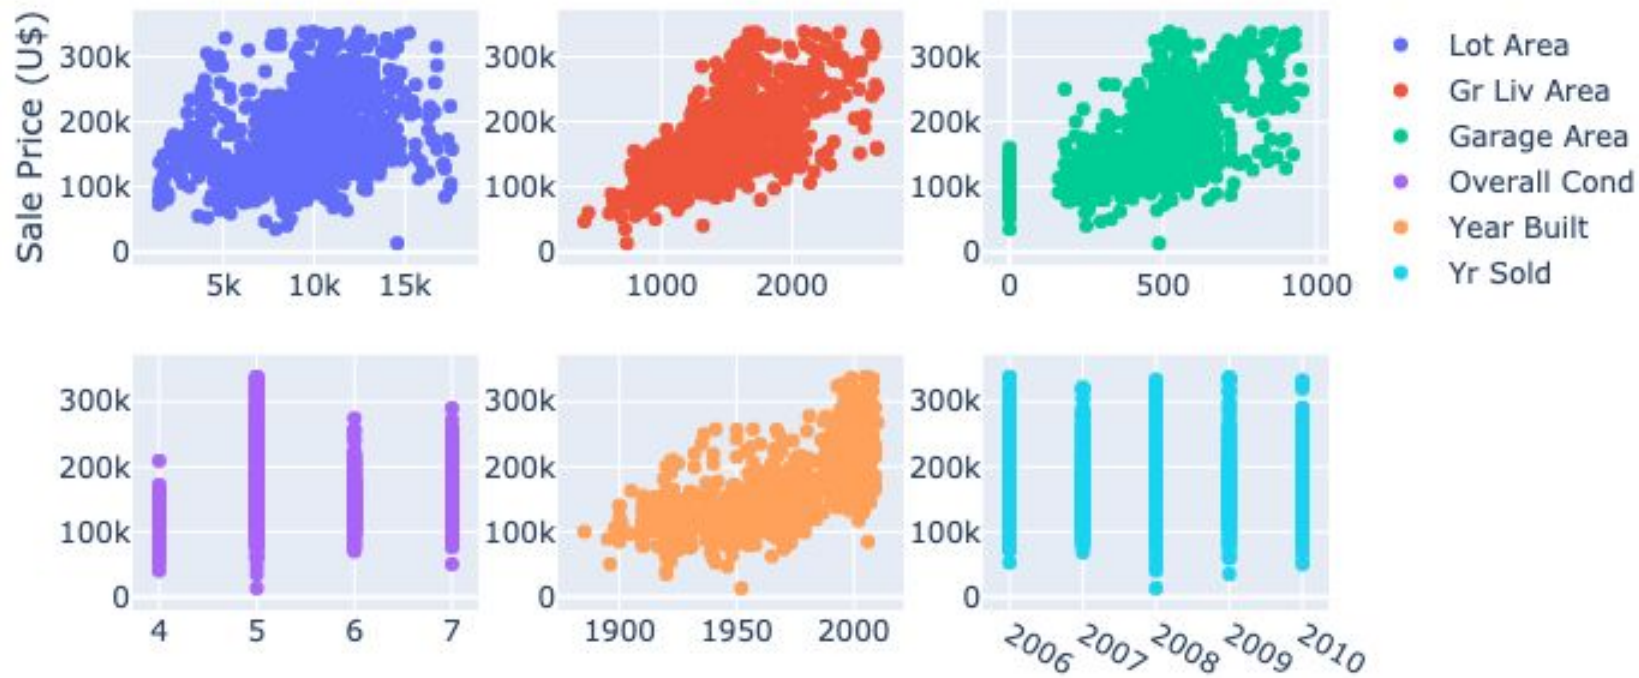
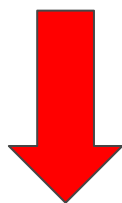
```
X_train, X_test, Y_train, Y_test = train_test_split(data_iqr.drop(axis=1, labels=[ "SalePrice" ]),  
                                                    data_iqr[ "SalePrice" ],  
                                                    test_size=0.5,  
                                                    random_state=42)
```



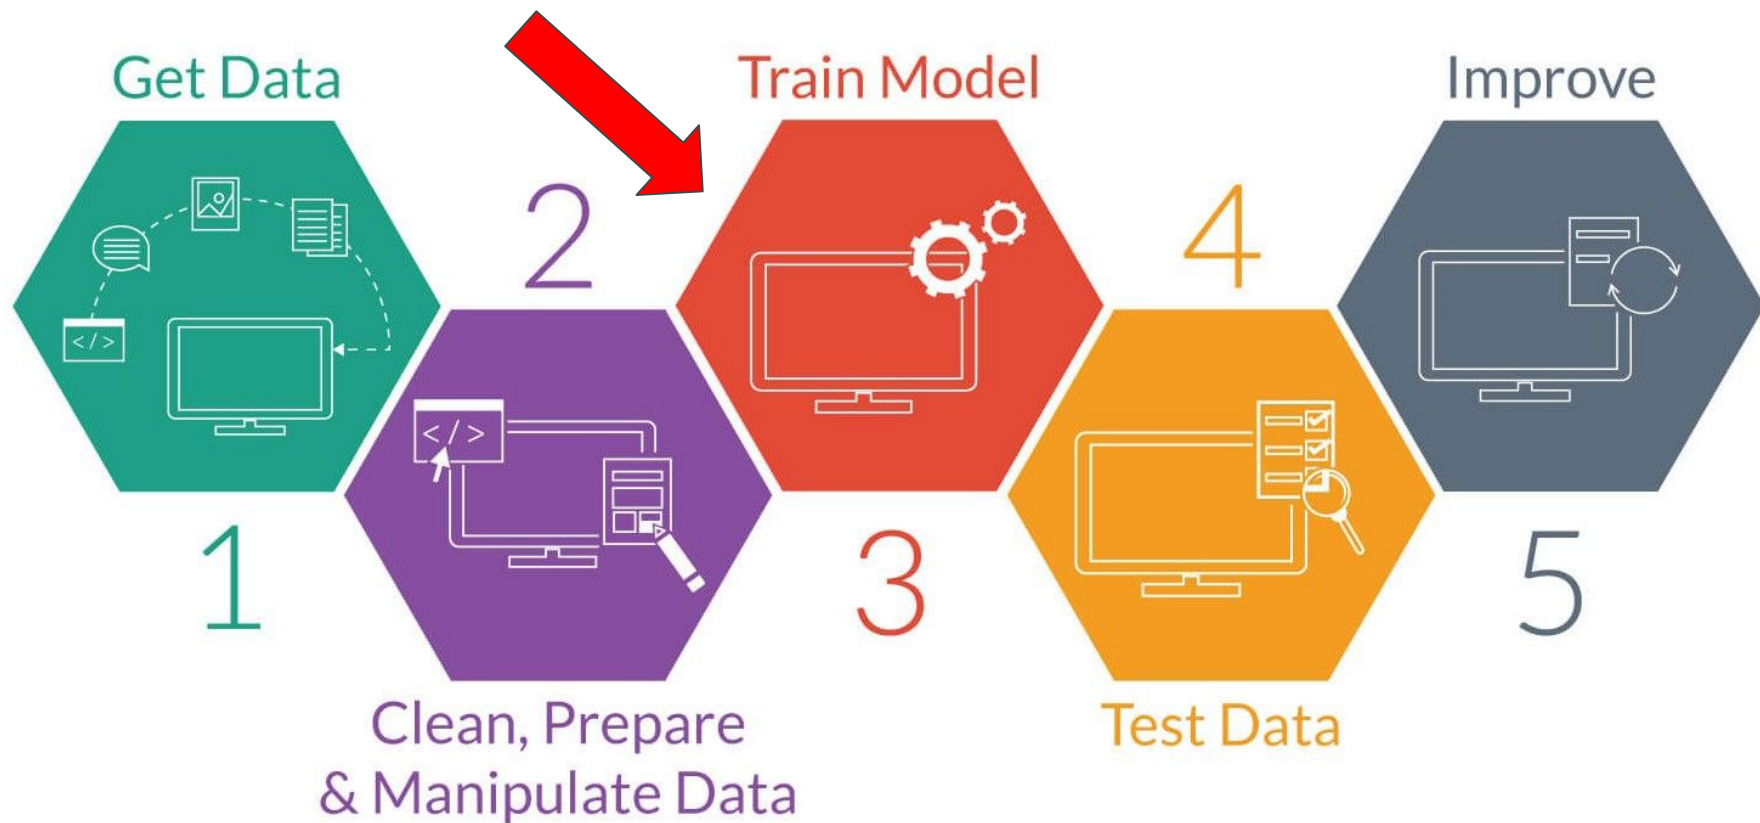
```
pd.concat([X_train,Y_train],axis=1).corr()["SalePrice"].sort_values()
```



| | |
|--------------|-----------|
| Overall Cond | -0.186210 |
| Yr Sold | 0.017669 |
| Lot Area | 0.305555 |
| Garage Area | 0.632527 |
| Year Built | 0.650917 |
| Gr Liv Area | 0.676906 |
| SalePrice | 1.000000 |



A general ML workflow



Linear Regression with **One Variable**

Notation:

- m - number of training examples
- X 's - input variable/features
- y 's - output variable/ target variable

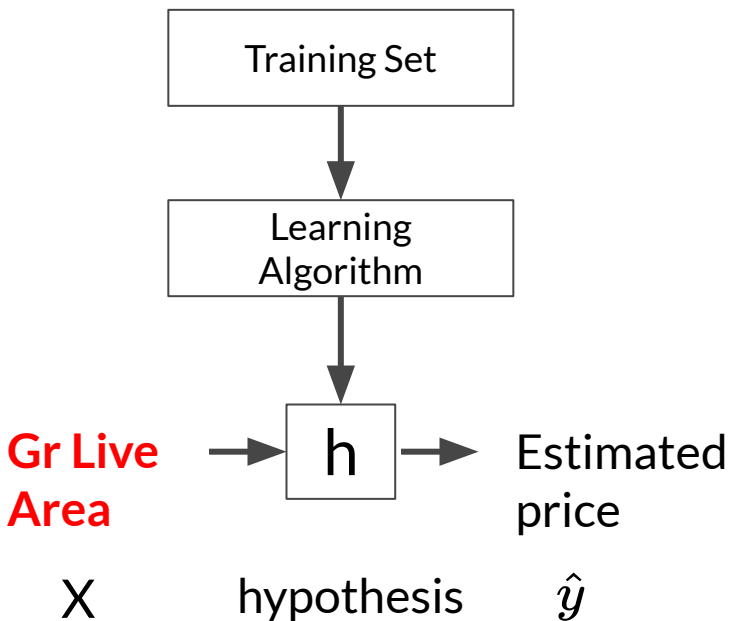
$$\begin{array}{ll} X^{(1)} = 1173 & y^{(1)} = 170000 \\ X^{(2)} = 1096 & y^{(2)} = 138800 \\ X^{(3)} = 1012 & y^{(3)} = 127500 \end{array}$$

$(X^{(i)}, y^{(i)}) = i^{\text{th}}$ training example

$m = 1202$

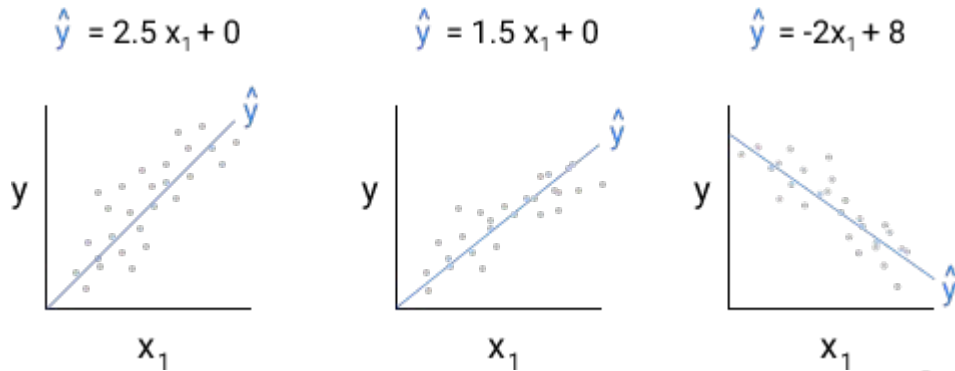
| | | X | y |
|----------|--|-------------|-----------|
| | | Gr Liv Area | SalePrice |
| m = 1202 | | 1173 | 170000 |
| | | 1096 | 138800 |
| | | 1012 | 127500 |
| | | 1797 | 231000 |
| | | 1436 | 225000 |

Model Representation (Linear Reg. **One Variable**)



How do we represent h ?

$$\hat{y} = h_{\theta}(x) = \theta_1 x_1 + \theta_0$$



Cost Function

$$f(x)$$

"minimize the error"

Cost Function (Linear Reg. One Var.)

Hypothesis $h_{\theta}(x) = \theta_0 + \theta_1 x$

θ_i = parameters

How to choose θ_i ?

m = 1202

| X | y |
|-------------|-----------|
| Gr Liv Area | SalePrice |
| 1173 | 170000 |
| 1096 | 138800 |
| 1012 | 127500 |
| 1797 | 231000 |
| 1436 | 225000 |

Cost Function

Intuition #01

(Linear Reg. One Var.)

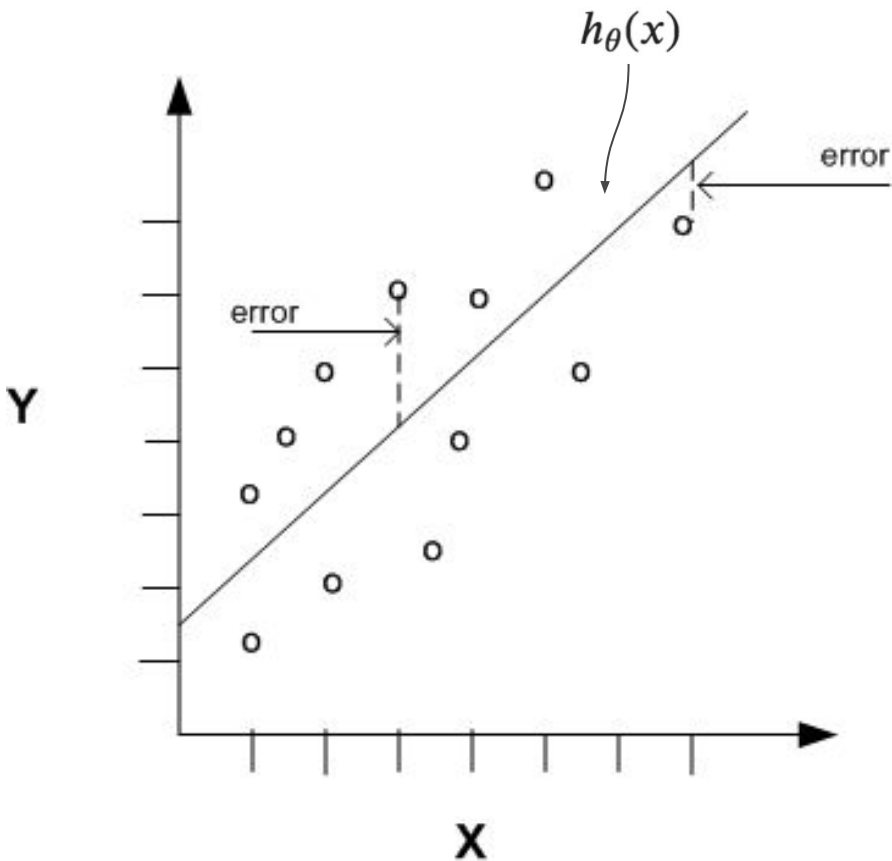
Cost Function

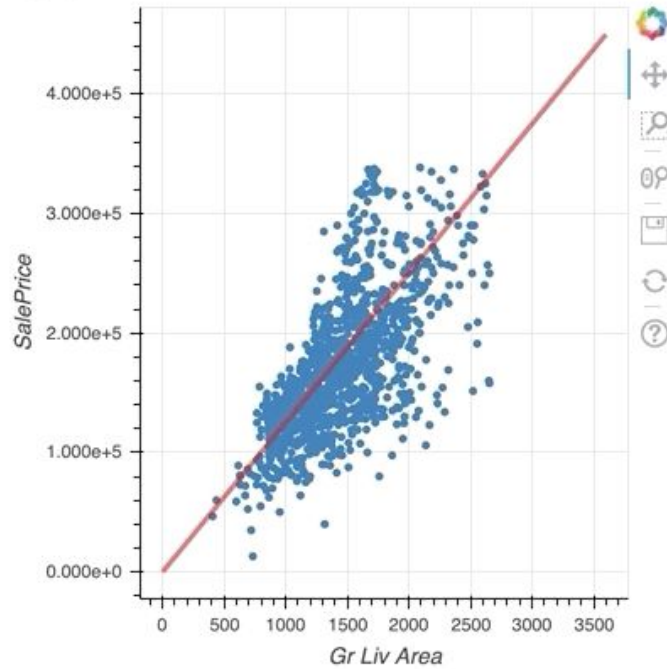
$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}]^2$$

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

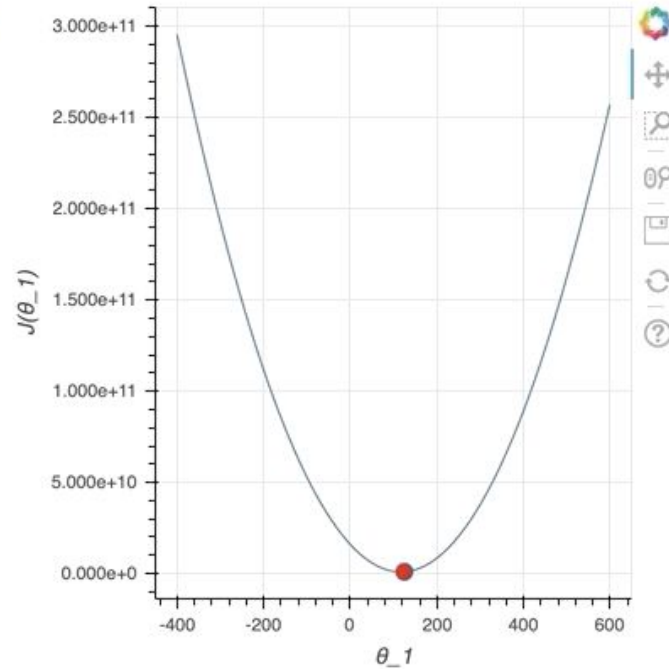
Idea:

- choose θ_0, θ_1 so that $h_{\theta}(x)$ is close to y for our training examples $(x^{(i)}, y^{(i)})$
- minimize (θ_0, θ_1)





A1: 125



$$\hat{y} = h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$\hat{y} = h_{\theta}(x) = \theta_1 x$$

Cost Function **Intuition #01**
 $(\theta_0 = 0)$

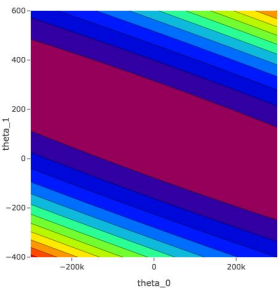
$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m [\theta_1 x^{(i)} - y^{(i)}]^2$$

Cost Function

Intuition #02

(Linear Reg. One var)

Contour



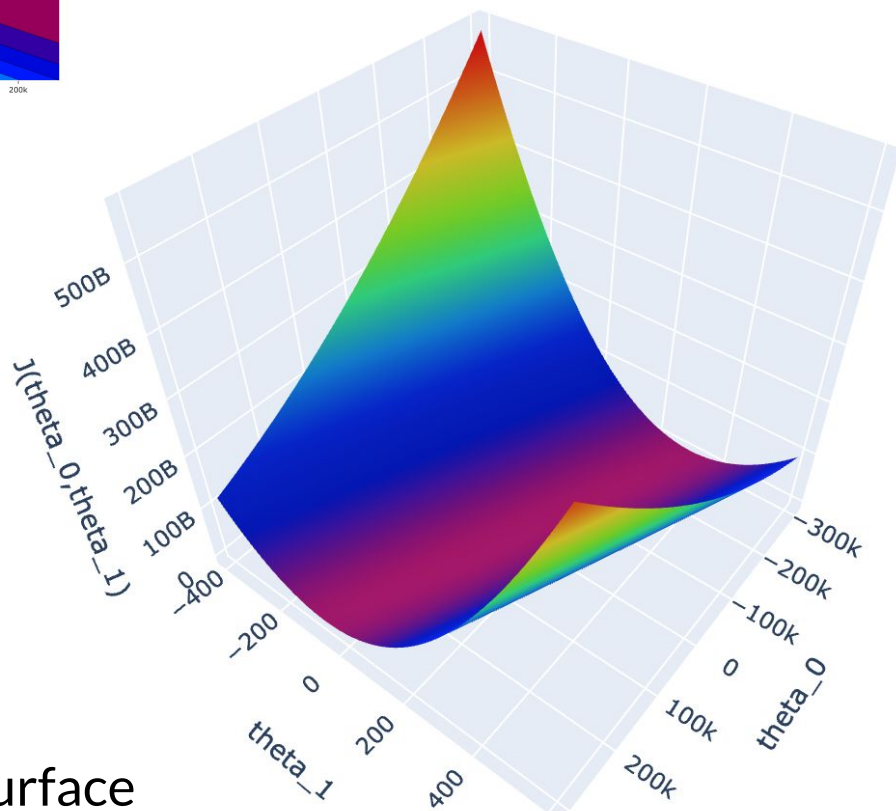
Cost Function

Intuition #02

(θ_0 and θ_1 are defined)

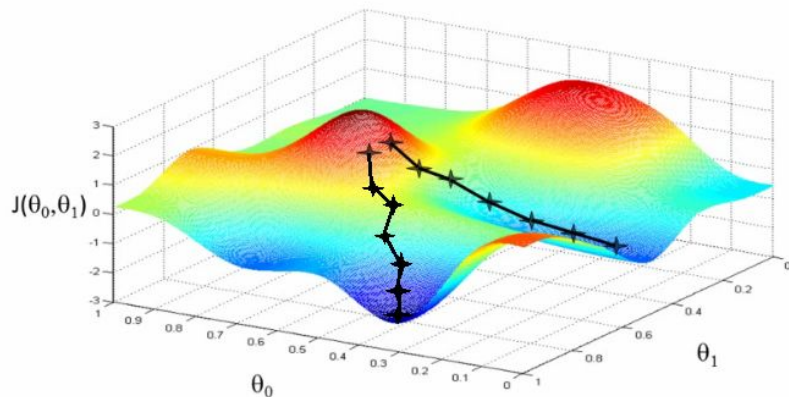
$$\hat{y} = h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}]^2$$

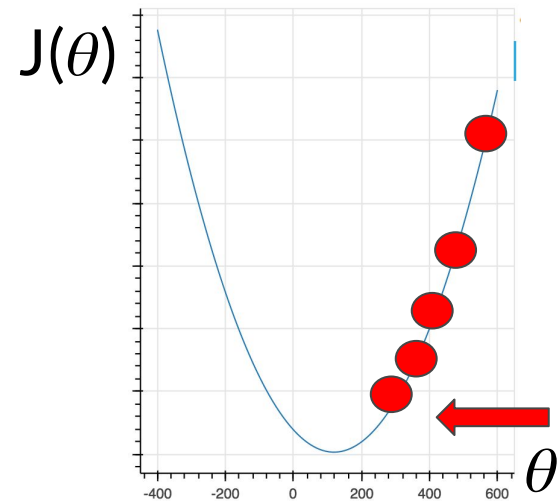


Surface

Gradient Descent (Linear Reg. One var)



1. Iteratively
 - 1.1. Evaluate parameters
 - 1.2. Compute loss
 - 1.3. Take small steps in the direction that will minimize loss

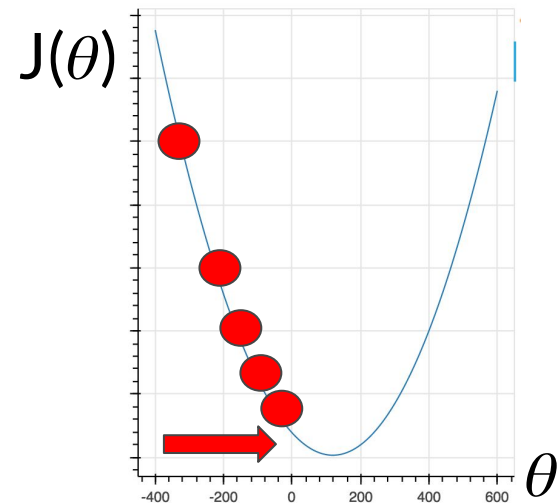


repeat until converge {

$$\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

}

α - learning rate



repeat until converge {

$$\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

}

Correct update

$$aux_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$aux_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_0 = aux_0$$

$$\theta_1 = aux_1$$

Incorrect update

$$aux_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$\theta_0 = aux_0$$

$$aux_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_1 = aux_1$$

repeat until converge {

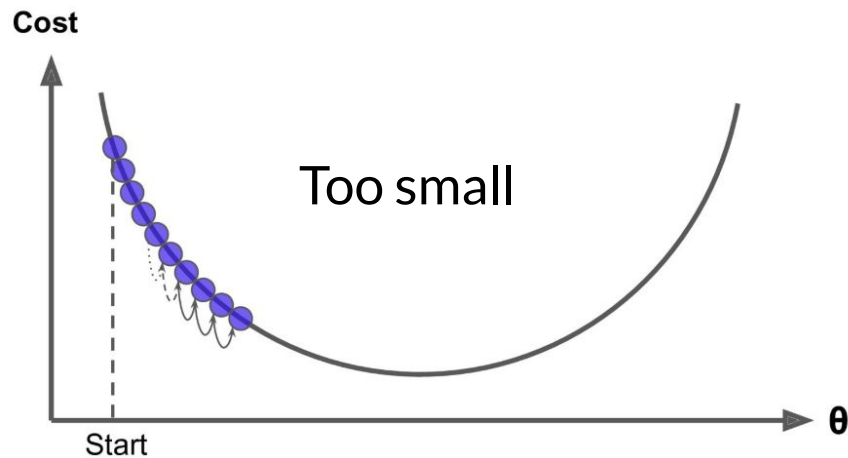
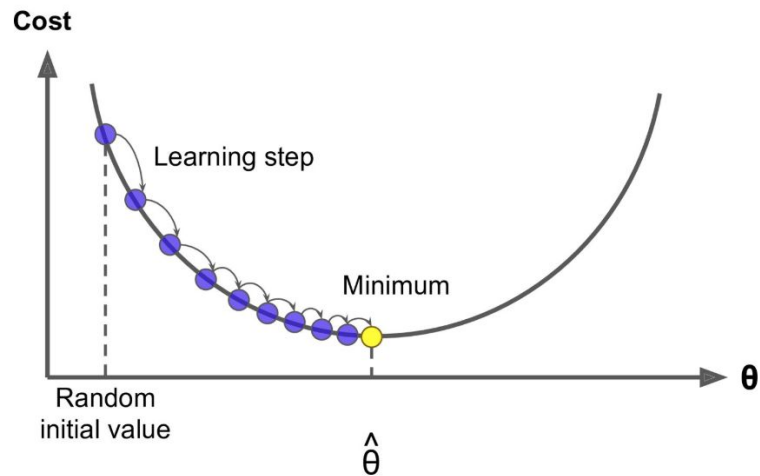
$$aux_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}]$$

$$aux_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}] x^{(i)}$$

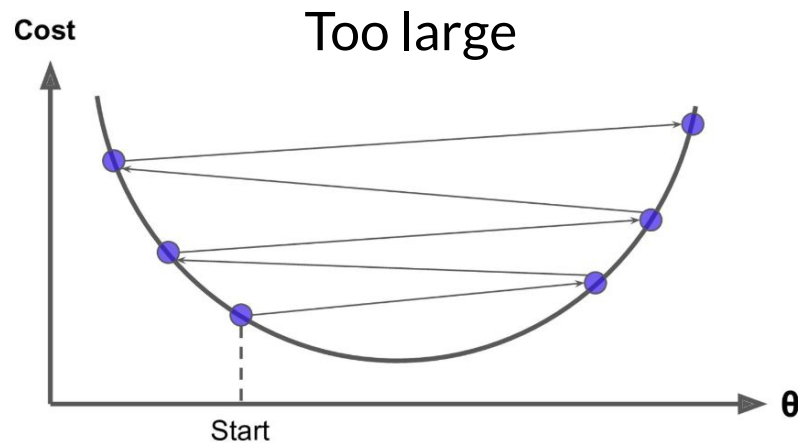
$$\theta_0 = aux_0$$

$$\theta_1 = aux_1$$

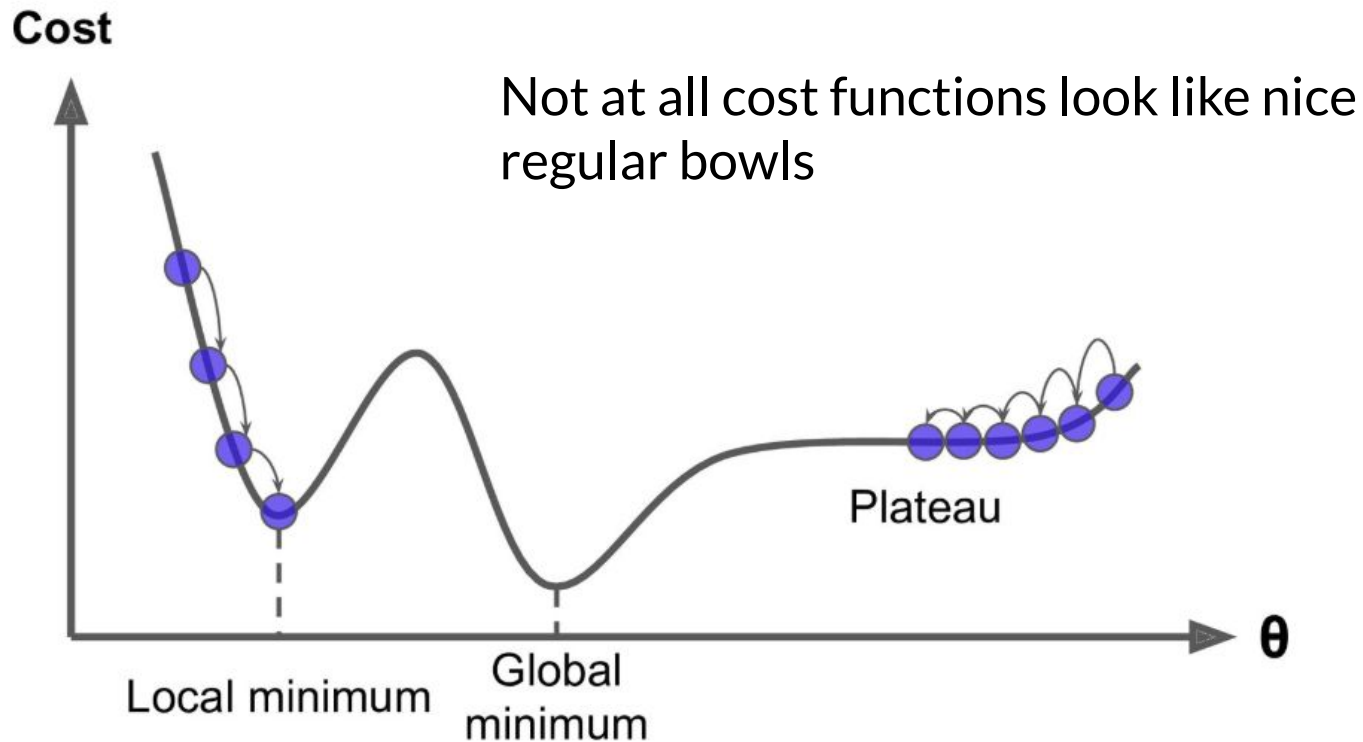
}



Learning rate
tradeoff



Gradient Descent Pitfalls



cost function & gradient descent from linear algebra perspective

Hypothesis

$$\hat{y} = h_{\theta}(x) = \theta_0 + \theta_1 x$$

| Gr Liv Area | SalePrice |
|-------------|-----------|
| 2480 | 205000 |
| 1829 | 237000 |
| 2673 | 249000 |
| 1005 | 133500 |
| 1768 | 224900 |

[Export to plot.ly »](#)

$$\text{hypothesis} = \begin{bmatrix} 1 & 2480 \\ 1 & 1829 \\ 1 & 2679 \\ 1 & 1005 \\ 1 & 1768 \end{bmatrix} \times \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix} = \begin{bmatrix} 2480 \theta_1 + \theta_0 \\ 1829 \theta_1 + \theta_0 \\ 2679 \theta_1 + \theta_0 \\ 1005 \theta_1 + \theta_0 \\ 1768 \theta_1 + \theta_0 \end{bmatrix}$$

```
def cost_function(X, y, theta):
    return np.sum(np.square(np.matmul(X, theta) - y)) / (2 * len(y))
```

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}]^2$$

| Gr Liv Area | SalePrice |
|-------------|-----------|
| 2480 | 205000 |
| 1829 | 237000 |
| 2673 | 249000 |
| 1005 | 133500 |
| 1768 | 224900 |

[Export to plot.ly »](#)

$$J(\theta_0, \theta_1) = \frac{1}{2 \times 5} \sum \left(\begin{bmatrix} 2480 \theta_1 + \theta_0 \\ 1829 \theta_1 + \theta_0 \\ 2679 \theta_1 + \theta_0 \\ 1005 \theta_1 + \theta_0 \\ 1768 \theta_1 + \theta_0 \end{bmatrix} - \begin{bmatrix} 205000 \\ 237000 \\ 249000 \\ 133500 \\ 224900 \end{bmatrix} \right)^2$$

repeat until converge {

$$aux_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}]$$

$$aux_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m [h_{\theta}(x^{(i)}) - y^{(i)}] x^{(i)}$$

$$\theta_0 = aux_0$$

$$\theta_1 = aux_1$$

}

```
def gradient_descent(X, y, alpha, iterations, theta):
```

```
    m = len(y)
```

```
    cost_history = []
```

```
    for i in range(iterations):
```

```
        t0 = theta[0] - (alpha / m) * np.sum(np.dot(X, theta) - y)
```

```
        t1 = theta[1] - (alpha / m) * np.sum((np.dot(X, theta) - y) * X[:,1])
```

```
        theta = np.array([t0, t1])
```

```
        cost_history.append(cost_function(X, y, theta))
```

```
    return theta, cost_history
```

- Involves calculations over the full training set X at each gradient step
- It uses the whole batch of training data at every step.
- This is why the algorithm called **Batch Gradient Descent**

1. Batch gradient descent: use all **m examples** in each iteration
2. Stochastic gradient descent: use **1 example** in each iteration
3. Mini-batch gradient descent: use **b examples** in each iteration

Stochastic gradient descent

Randomly shuffle (reorder)
training examples

Repeat {

 for $i := 1, \dots, m$ {

$$\theta_j := \theta_j - \alpha(h_{\theta}(x^{(i)}) - y^{(i)})x_j^{(i)}$$

 (for every $j = 0, \dots, n$)

 }

}

Stochastic gradient descent

```
from sklearn.utils import shuffle
```

```
# stochastic gradient descent
```

```
def stochastic_gradient_descent(X, y, alpha, iterations, theta):
```

```
    m = len(y)
```

```
    cost_history = []
```

```
    # randomly shuffle the training dataset
```

```
    X,y = shuffle(X,y,random_state=42)
```

```
    for i in range(iterations):
```

```
        for j in range(m):
```

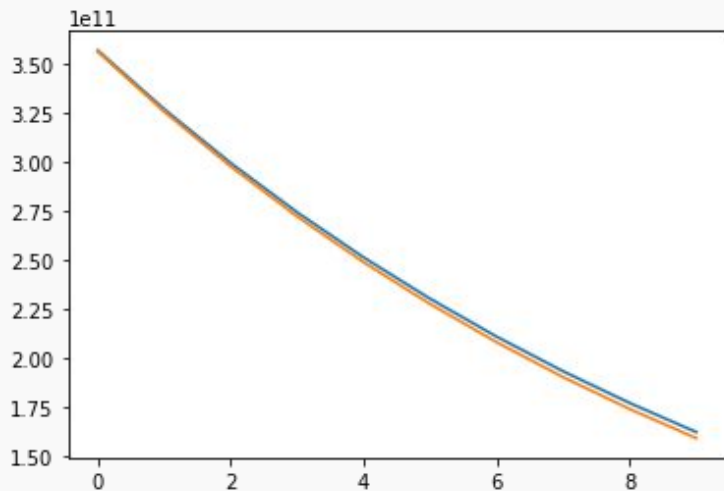
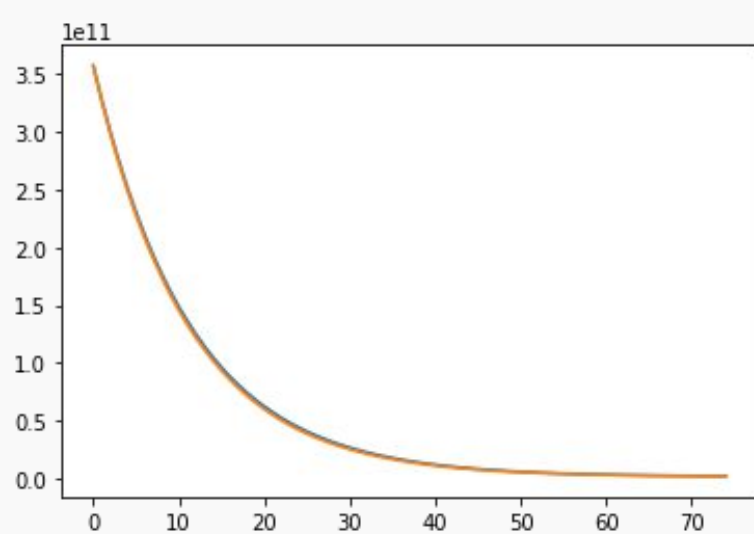
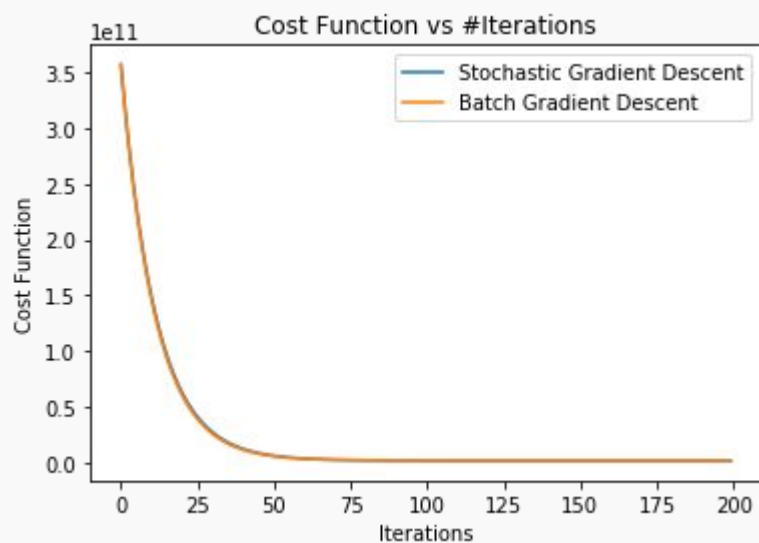
```
            t0 = theta[0] - alpha * np.sum(np.dot(X[j,:], theta) - y[j])
```

```
            t1 = theta[1] - (alpha / m) * np.sum((np.dot(X[j,:], theta) - y[j]) * X[j,1])
```

```
            theta = np.array([t0, t1])
```

```
            cost_history.append(cost_function(X, y, theta))
```

```
    return theta, cost_history
```



Stochastic Gradient Descent

- Pro: faster learning, can avoid local minima
- Cons: computationally expensive

Batch Gradient Descent

- Pro: computationally efficient, stable convergence
- Cons: memory--

Mini-Batch gradient descent

Say $b = 10, m = 1000$.

Repeat {

for $i = 1, 11, 21, 31, \dots, 991$ {

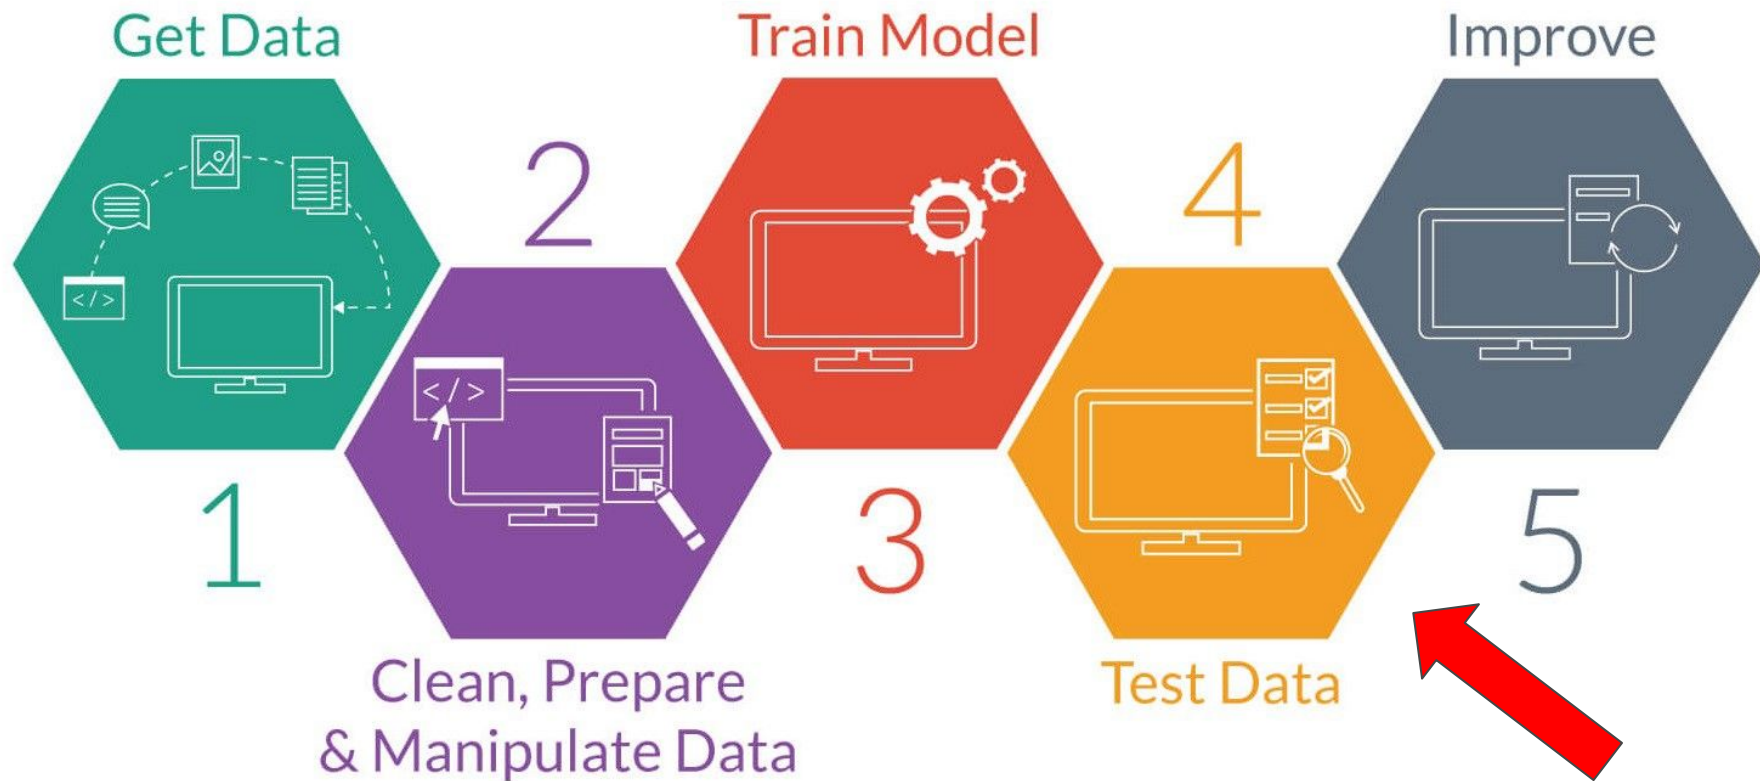
$$\theta_j := \theta_j - \alpha \frac{1}{10} \sum_{k=i}^{i+9} (h_{\theta}(x^{(k)}) - y^{(k)}) x_j^{(k)}$$

(for every $j = 0, \dots, n$)

}

}

A general ML workflow



The Linear Regression Model.ipynb

Section #01





- We are going to start by covering linear regression
 - Multiple variables
- We discuss the application of linear regression to **housing price prediction**

Linear Regression with Multiple Variables

Notation:

- m - number of training examples
- n - number of features
- $x^{(i)}$ - input features of i^{th} training example
- $x_j^{(i)}$ - value of feature j in i^{th} training example
- $y^{(i)}$ - target value of i^{th} training examples

$$x^{(2)} = \begin{bmatrix} 11622 \\ 5 \\ 1961 \\ 2010 \\ 105000 \end{bmatrix}$$

$$x_3^{(2)} = 1961$$

$n = 4$

| x_1 | x_2 | x_3 | x_4 | y |
|----------|--------------|------------|---------|-----------|
| Lot Area | Overall Qual | Year Built | Yr Sold | SalePrice |
| 31770 | 6 | 1960 | 2010 | 215000 |
| 11622 | 5 | 1961 | 2010 | 105000 |
| 14267 | 6 | 1958 | 2010 | 172000 |
| 11160 | 7 | 1968 | 2010 | 244000 |
| 13830 | 5 | 1997 | 2010 | 189900 |

$m = 5$

Hypothesis (previously)

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Multivariable case

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4$$

For convenience of notation, define $x_0=1$. In other words:

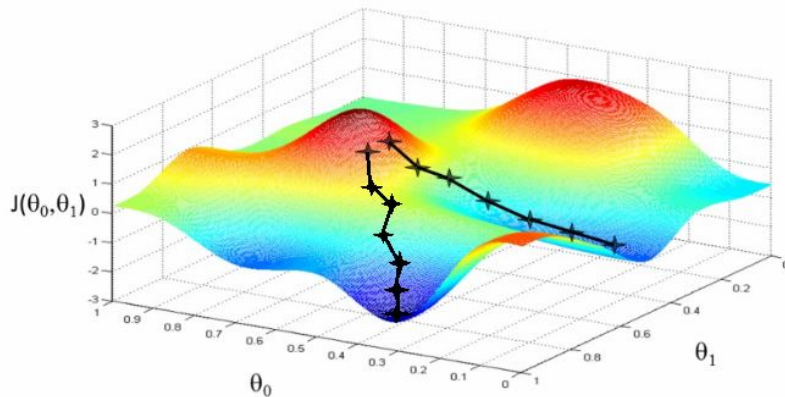
$$x_0^{(i)}=1$$

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1} \qquad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$h_{\theta}(x) = \begin{bmatrix} \theta_0 & \theta_1 & \theta_1 & \dots & \theta_n \end{bmatrix} \times \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = \theta^T x$$

Gradient Descent (Linear Reg. **Multiple Variables**)



Hypothesis: $h_{\theta}(x) = \theta^T x = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$

Parameters: $\theta_0, \theta_1, \theta_2, \dots, \theta_n$

Cost function:

$$J(\theta_0, \theta_1, \theta_2, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Gradient descent:

repeat {

$$\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \dots, \theta_n)$$

}

Simultaneously update for every j (0 to n)

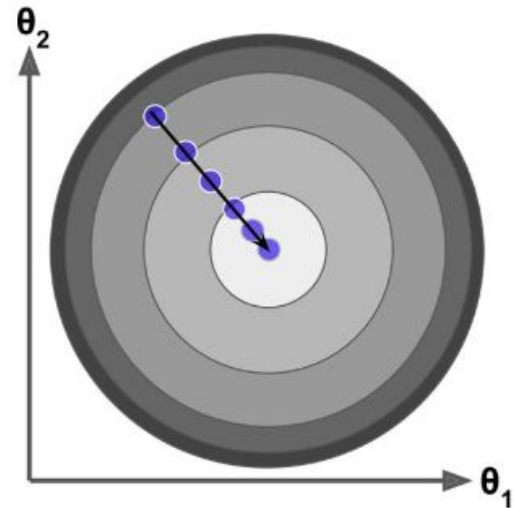
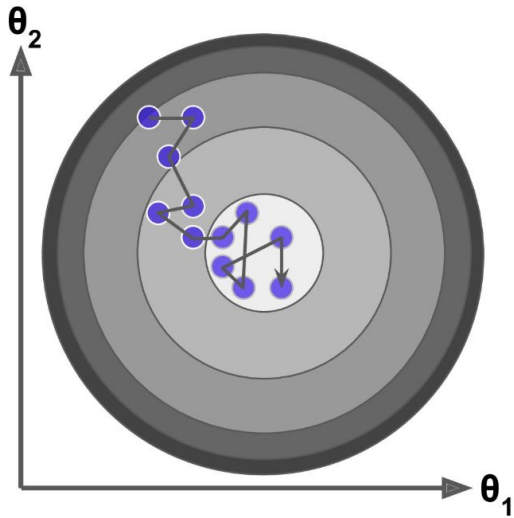
Gradient descent:

repeat until the convergence {

$$\theta_j = \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \quad \text{for } j = 0, \dots, n$$

}

Gradient Descent: **trick #1** - Feature Scaling



Gradient Descent: **trick #1** - Feature Scaling

Z-Score or Standardization

$$z = \frac{x - \mu}{\sigma}$$

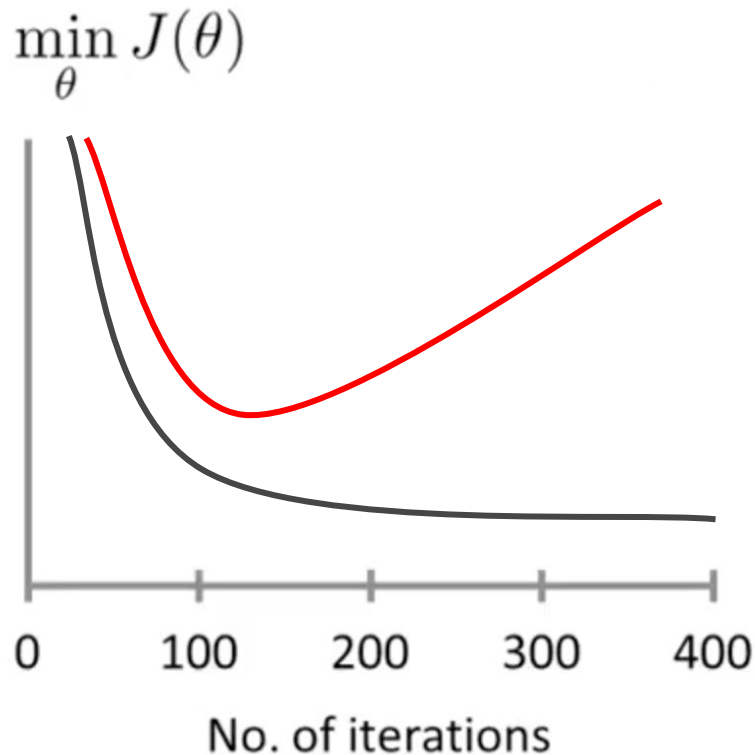
$\mu - 0$
 $\sigma - 1$

Min-Max Scaling

$$X_{norm} = \frac{X - X_{min}}{X_{max} - X_{min}}$$

range (0,1)

Gradient Descent: **trick #2** - Debugging α



1. Make a plot with number of iterations on the x-axis.
2. Now plot the cost function, $J(\theta)$ over the number of iterations of gradient descent.
3. If $J(\theta)$ ever increases, then you probably need to decrease α .
4. It has been proven that if learning rate α is sufficiently small, then $J(\theta)$ will decrease on every iteration.
5. Automatic convergence test

Gradient Descent: **trick #2** - Debugging α

- If α is too small:
 - Slow convergence
- If α is too large:
 - $J(\theta)$ may not decrease on every iteration;
 - $J(\theta)$ may not converge.

To choose α :

..., 0.0001, ..., 0.001, ..., 0.01, ..., 0.1, ..., 1, ..., 10, ...

normal equation: method to
solve for θ analytically

Normal Equation

| Lot Area | Overall Qual | Year Built | Yr Sold | SalePrice |
|----------|--------------|------------|---------|-----------|
| 31770 | 6 | 1960 | 2010 | 215000 |
| 11622 | 5 | 1961 | 2010 | 105000 |
| 14267 | 6 | 1958 | 2010 | 172000 |
| 11160 | 7 | 1968 | 2010 | 244000 |
| 13830 | 5 | 1997 | 2010 | 189900 |

$$X\theta = y$$

$$X^T X \theta = X^T y$$

$$\theta = (X^T X)^{-1} X^T y$$

$$\begin{bmatrix} 1 & 31770 & 6 & 1960 & 2010 \\ 1 & 11622 & 5 & 1961 & 2010 \\ 1 & 14267 & 6 & 1958 & 2010 \\ 1 & 11160 & 7 & 1968 & 2010 \\ 1 & 13830 & 5 & 1997 & 2010 \end{bmatrix} \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \theta_3 \\ \theta_4 \\ \theta_5 \end{bmatrix} = \begin{bmatrix} \theta_0 + 31770\theta_1 + 6\theta_2 + 1960\theta_3 + 2010\theta_4 \\ \theta_0 + 11622\theta_1 + 5\theta_2 + 1961\theta_3 + 2010\theta_4 \\ \theta_0 + 14267\theta_1 + 6\theta_2 + 1958\theta_3 + 2010\theta_4 \\ \theta_0 + 11160\theta_1 + 7\theta_2 + 1968\theta_3 + 2010\theta_4 \\ \theta_0 + 13830\theta_1 + 5\theta_2 + 1997\theta_3 + 2010\theta_4 \end{bmatrix}$$

m training examples, n features

Gradient Descent

- Need to choose α
- Needs many iterations
- Works well even when n is large

Normal Equation

- No need to choose α
- Don't need to iterate
- Need to compute $(X^T X)^{-1}$
- Slow if n is very large

If $(X^T X)$ is noninvertible, the common causes might be having:

- Redundant features, where two features are very closely related (i.e. they are linearly dependent)
- Too many features (e.g. $m \leq n$). In this case, delete some features or use "regularization" (to be explained in a later lesson)

